

Artificial Intelligence

Constraint Satisfaction Problems

Dr. Bilgin Avenoğlu

Constraint Satisfaction Problems

- Up to now, we explored the idea that problems can be solved by searching the state space
- Each state is atomic, or indivisible— a black box with no internal structure.

How the Components of Agent Programs Work?

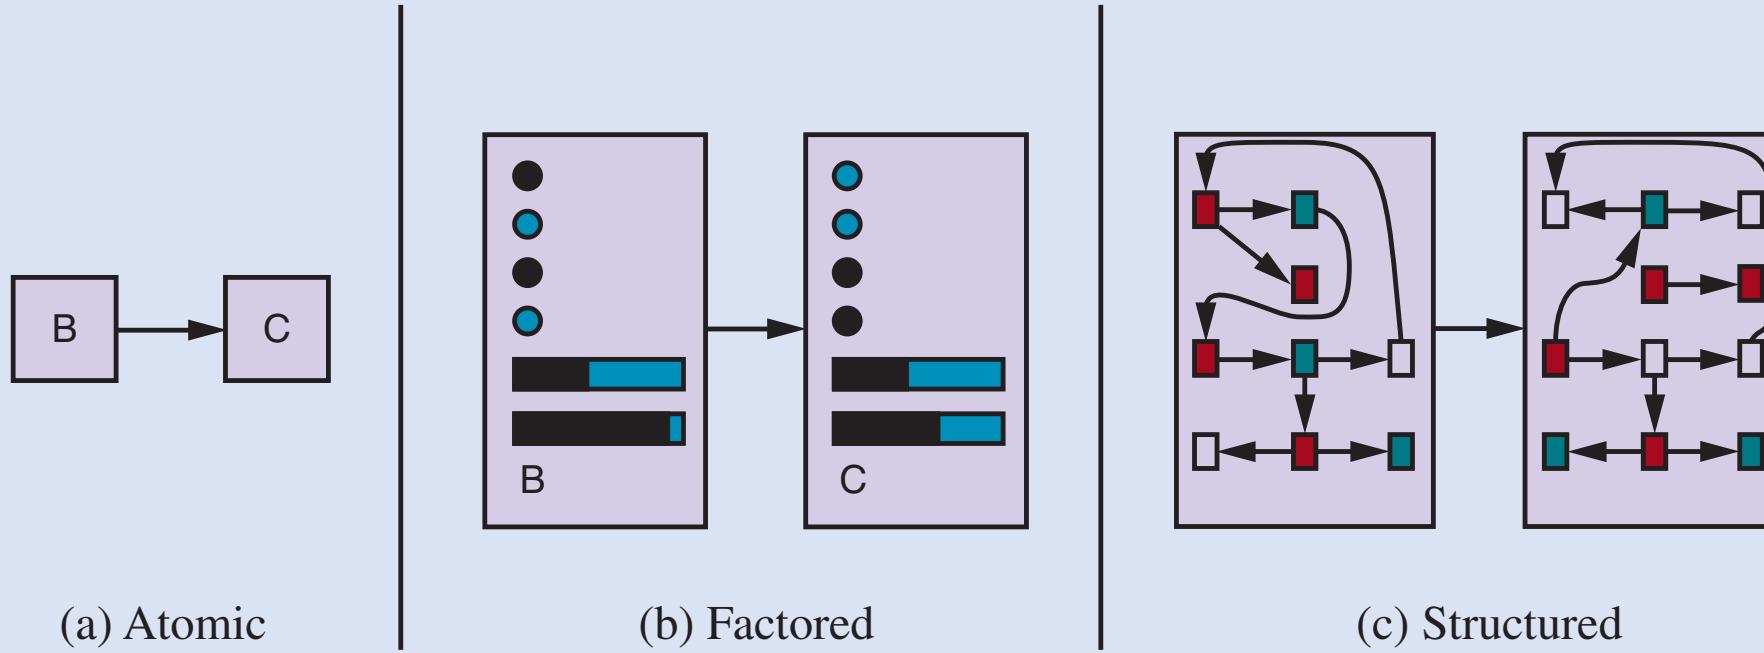


Figure 2.16 Three ways to represent states and the transitions between them. (a) Atomic representation: a state (such as B or C) is a black box with no internal structure; (b) Factored representation: a state consists of a vector of attribute values; values can be Boolean, real-valued, or one of a fixed set of symbols. (c) Structured representation: a state includes objects, each of which may have attributes of its own as well as relationships to other objects.

Constraint Satisfaction Problems

- Now, we **break** open the **black box** by using a factored representation for each state:
 - a **set of variables**, each of which **has** a **value**.
- A problem is solved when **each variable** has a value that **satisfies** all the **constraints** on the variable.
- A problem described this way is called a **constraint satisfaction problem - CSP**.

Constraint Satisfaction Problems

- **CSP** search algorithms **use general rather than domain-specific heuristics** to enable the solution of complex problems.
 - The main idea is to **eliminate large portions of the search space all at once** by identifying variable/value combinations that **violate the constraints**.
- **CSPs** have the additional **advantage** that the **actions** and **transition model** can be **deduced from the problem description**.

Defining Constraint Satisfaction Problems

A constraint satisfaction problem consists of three components, \mathcal{X} , \mathcal{D} , and \mathcal{C} :

\mathcal{X} is a set of variables, $\{X_1, \dots, X_n\}$.

\mathcal{D} is a set of domains, $\{D_1, \dots, D_n\}$, one for each variable.

\mathcal{C} is a set of constraints that specify allowable combinations of values.

- if X_1 and X_2 both have the domain $\{1,2,3\}$, then the constraint saying that X_1 must be greater than X_2 can be written as
 - $\langle(X_1, X_2), \{(3,1),(3,2),(2,1)\}\rangle$ or as $\langle(X_1, X_2), X_1 > X_2 \rangle$.

Defining Constraint Satisfaction Problems

- CSPs deal with **assignments** of values to variables,
 - $\{X_i = v_i, X_j = v_j, \dots\}$.
- An assignment that does **not violate any constraints** is called a **consistent** or legal assignment.
- A **complete** assignment is one in which **every variable** is **assigned** a **value**,
- A **solution** to a CSP is a **consistent, complete** assignment.
- A **partial assignment** is one that **leaves some** variables **unassigned**, and a partial solution is a partial assignment that is consistent.

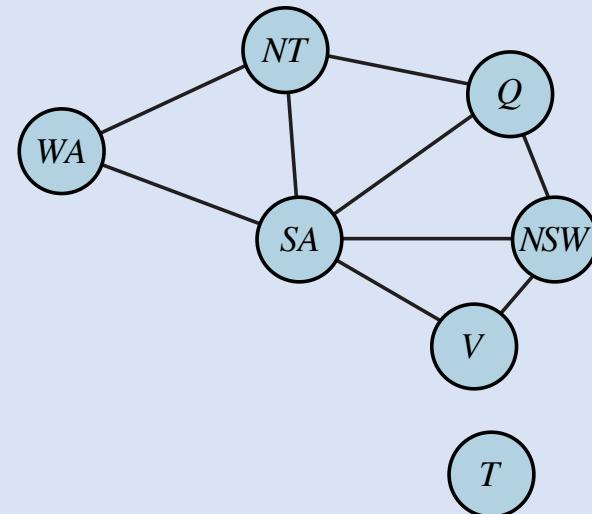
Example problem: Map coloring

- Coloring each region either red, green, or blue
 - no two neighboring regions have the same color.
 - $X = \{WA, NT, Q, NSW, V, SA, T\}$.
- The **domain** of every variable is the set $D_i = \{\text{red, green, blue}\}$.
- The constraints require neighboring regions to have distinct colors.
 - $C = \{SA \neq WA, SA \neq NT, SA \neq Q, SA \neq NSW, SA \neq V, WA \neq NT, NT \neq Q, Q \neq NSW, NSW \neq V\}$.
- $SA \neq WA$ is a shortcut for $\langle (SA, WA), SA \neq WA \rangle$, where $SA \neq WA$ can be fully enumerated in turn as
 - $\{(red, green), (red, blue), (green, red), (green, blue), (blue, red), (blue, green)\}$.
- There are many **possible solutions** to this problem, such as
 - $\{WA=\text{red}, NT=\text{green}, Q=\text{red}, NSW=\text{green}, V=\text{red}, SA=\text{blue}, T=\text{red}\}$.



Constraint graph

- The **nodes** of the graph **correspond** to **variables** of the problem, and an **edge** connects any two variables that participate in a **constraint**.



Why formulate a problem as a CSP?

- It is often **easy to formulate** a problem as a CSP
- CSP **solvers fast** and **efficient**.
- A CSP **solver** can quickly **prune** large swathes of the **search space** that an **atomic state-space searcher** cannot.
- In atomic state-space search we can only ask: is this specific state a goal? No? What about this one?
 - With CSPs, once we **find out** that a **partial assignment violates a constraint**, we can **immediately discard** further refinements of the partial assignment.

Types of variables

- CSP involves variables that have **discrete, finite** domains.
 - Map-coloring problem is of this kind.
 - The **8-queens** problem can also be viewed as a **finite-domain** CSP,
 - the variables Q_1, \dots, Q_8 correspond to the queens in columns 1 to 8,
 - the domain of each variable specifies the possible row numbers for the queen in that column,
 $D_i = \{1, 2, 3, 4, 5, 6, 7, 8\}$.
 - The **constraints** say that no two queens can be in the same row or diagonal.

Types of constraints

- A *unary constraint* : restricts the value of a single variable.
 - South Australians won't tolerate the color green; $\langle (SA), SA \neq \text{green} \rangle$.
- A *binary constraint* relates two variables.
 - $SA \neq NSW$ is a binary constraint.
- A *binary CSP* is one with **only unary and binary** constraints; it can be represented as a **constraint graph**
- We can also define **higher-order** constraints.
 - The **ternary** constraint *Between*(X, Y, Z),
 - $\langle (X, Y, Z), X < Y < Z \text{ or } X > Y > Z \rangle$.

Types of constraints

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7								8
F		6	7	8	2				
G		2	6	9	5				
H	8		2	3				9	
I		5	1		3				

- A constraint involving an **arbitrary number of variables** is called a *global constraint*.
- A common global constraints: *Alldiff*
 - All of the variables involved in the constraint must have different values.
 - In **Sudoku** problems, all **variables in a row, column, or 3×3 box** must **satisfy** an *Alldiff* constraint.
- Every finite-domain constraint can be reduced to a set of binary constraints if enough auxiliary variables are introduced.
- We can transform any **CSP** into one with only binary constraints
 - certainly makes the life of the algorithm designer simpler.

Constraint Propagation: Inference in CSPs

- **Node consistency**

- A single variable (a node in the CSP graph) is **node-consistent** if all the values in the variable's domain **satisfy** the variable's **unary constraints**.
 - South Australians **dislike green**, the variable *SA* starts with domain $\{red, green, blue\}$, and we can make it **node consistent** by eliminating **green**, leaving SA with the reduced domain $\{red, blue\}$.
- We say that **a graph is node-consistent if every variable in the graph is node-consistent**.

Arc consistency

- A variable in a CSP is **arc-consistent** if every value in its domain **satisfies** the variable's **binary constraints**.
- X_i is **arc-consistent** with respect to another variable X_j if for **every value** in the current domain D_i , there is some value in the domain D_j that satisfies the binary constraint on the arc (X_i, X_j) .
- A graph is **arc-consistent** if every variable is arc-consistent with every other variable.
 - For example, consider the constraint $Y = X^2$ where the domain of both X and Y is the set of **decimal digits**. We can write this constraint explicitly as
 - $\langle (X, Y), \{(0, 0), (1, 1), (2, 4), (3, 9)\} \rangle$.
 - To make X **arc-consistent** with respect to Y , we **reduce** X 's domain to $\{0, 1, 2, 3\}$.
 - If we also make Y **arc-consistent** with respect to X , then Y 's domain becomes $\{0, 1, 4, 9\}$, and the **whole CSP** is **arc-consistent**.

Arc consistency example

- For $x \in [2 \dots 6]$, $y \in [3 \dots 7]$ the constraint

$$x < y$$

is arc-consistent.

- For example
 - if $x = 2$ then there is a solution
 - if $x = 6$ then y must be 7
 - if $y = 3$ then x must be 2.

- For $x \in [2..7]$ and $y \in [3..7]$ the constraint:

$$x < y$$

is not arc consistent.

- If $x = 7$ (allowed by the domains) then there is no value for y satisfying the constraint.

AC-3 algorithm

- AC-3 maintains a queue of arcs to consider.
- Initially, the queue contains all the arcs in the CSP.
 - Each binary constraint becomes two arcs, one in each direction.
- AC-3 then pops off an arbitrary arc (X_i, X_j) from the queue and makes X_i arc-consistent with respect to X_j .
 - If this leaves D_i unchanged, the algorithm just moves on to the next arc.
 - If this revises D_i (makes the domain smaller), then we add to the queue all arcs (X_k, X_i) where X_k is a neighbor of X_i .
 - The change in D_i might enable further reductions in D_k , even if we have previously considered X_k .
 - If D_i is revised down to nothing, then we know the whole CSP has no consistent solution.
- Otherwise, we keep checking, trying to remove values from the domains of variables until no more arcs are in the queue.
- We are left with a CSP that is equivalent to the original CSP
 - they both have the same solutions, but the arc-consistent CSP will be faster to search because its variables have smaller domains.
- In some cases, it solves the problem completely (by reducing every domain to size 1) and in others it proves that no solution exists (by reducing some domain to size 0).

```
function AC-3(csp) returns false if an inconsistency is found and true otherwise
  queue  $\leftarrow$  a queue of arcs, initially all the arcs in csp

  while queue is not empty do
     $(X_i, X_j) \leftarrow \text{POP}(\textit{queue})$ 
    if REVISE(csp,  $X_i$ ,  $X_j$ ) then
      if size of  $D_i = 0$  then return false
      for each  $X_k$  in  $X_i.\text{NEIGHBORS} - \{X_j\}$  do
        add  $(X_k, X_i)$  to queue
  return true

function REVISE(csp,  $X_i$ ,  $X_j$ ) returns true iff we revise the domain of  $X_i$ 
  revised  $\leftarrow$  false
  for each  $x$  in  $D_i$  do
    if no value  $y$  in  $D_j$  allows  $(x,y)$  to satisfy the constraint between  $X_i$  and  $X_j$  then
      delete  $x$  from  $D_i$ 
      revised  $\leftarrow$  true
  return revised
```

Arc consistency (AC-3)

- Example

- <https://www.boristhebrave.com/2021/08/30/arc-consistency-explained/>

- Start with a worklist containing two arcs for each constraint.
- While the worklist is non-empty:
 - Remove an arc from the worklist, say from x to y .
 - For each value in the domain of y :
 - Search for a support of that value on the constraint of the arc
 - If no support is found:
 - Remove the value from the domain of the variable
 - Add arcs from y to any other variables it's constrained with (excluding x).

Complexity of AC-3

- Assume a CSP with n variables, each with domain size at most d , and with c binary constraints (arcs).
- Each arc (X_k, X_i) can be inserted in the queue only d times because X_i has at most d values to delete.
- For each value in one domain, REVISE method examines all the values of the second domain.
 - So the complexity of the REVISE method would be d^2 .
- REVISE is called at maximum $c + cd$ times.
- The complexity of the whole algorithm is $O((c+cd)d^2)$ which is $O(cd^3)$.

Path Consistency

idea of arc consistency:

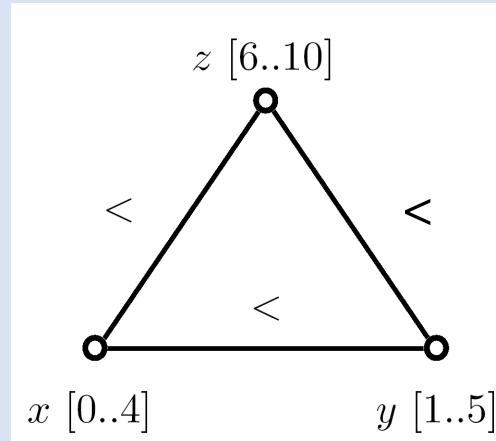
- For every assignment to a variable u there must be a suitable assignment to every other variable v .
- If not: remove values of u for which no suitable “partner” assignment to v exists.
~~~ tighter **unary constraint** on  $u$

This idea can be extended to three variables (**path consistency**):

- For every joint assignment to variables  $u, v$  there must be a suitable assignment to every third variable  $w$ .
- If not: remove pairs of values of  $u$  and  $v$  for which no suitable “partner” assignment to  $w$  exists.  
~~~ tighter **binary constraint** on  $u$  and  $v$

Example

$\langle x < y, y < z, x < z ; x \in [0..4], y \in [1..5], z \in [6..10] \rangle$ path consistent



$$C_{x,y} = \{(a,b) \mid a < b, a \in [0..4], b \in [1..5]\}$$

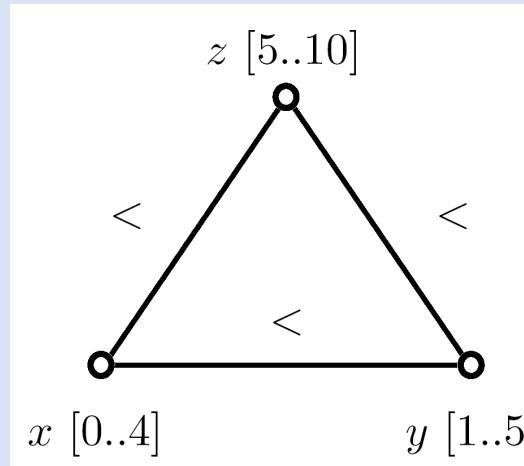
$$C_{x,z} = \{(a,c) \mid a < c, a \in [0..4], c \in [6..10]\}$$

$$C_{y,z} = \{(b,c) \mid b < c, b \in [1..5], c \in [6..10]\}$$

→ the 3 conditions (cf. previous slide) are satisfied

Example

$\langle x < y, y < z, x < z ; x \in [0..4], y \in [1..5], z \in [5..10] \rangle$ not path consistent



$$C_{x,z} = \{(a,c) \mid a < c, a \in [0..4], c \in [5..10]\}$$

But for $4 \in [0..4]$ and $5 \in [5..10]$ there is no $y \in [1..5]$ s.t. $4 < y$ and $y < 5$.

PC-6

- The **best path-consistency algorithm** proposed is PC-6, a natural generalization of AC-6 to path-consistency.
 - Its time complexity is $O(n^3d^3)$
 - Its space complexity is $O(n^3d^2)$,
 - where n is the number of variables and d is the size of domains.

Strong Consistency

Define P as a CSP where X, Y are the variables, domain of both is {1,2,3,4} and conditions in normal form are:

1. node-condition $X < 4$
2. arc-condition $X = Y$

P is 2-consistent (arc consistent) because for any X value it is possible to find a Y value that fulfills the arc-condition $X = Y$.

However, P is not 1-consistent (node-consistent) because exist a X value ($X = 4$) that can not fulfill the node condition $x < 4$.

For these reasons, this problem is 2-consistent but not strongly 2-consistent.

Obviously, it is straightforward to convert this example in a strongly 2-consistent problem, just reducing the domain to {1,2,3}.

K-consistency

- A CSP is **strongly k-consistent** if it is k-consistent and is also
 - $(k - 1)$ -consistent, $(k - 2)$ - consistent, ... all the way down to 1-consistent.
- Suppose we have a CSP with n nodes and make it **strongly n-consistent**.
 - First, we **choose a consistent value** for X_1 .
 - We are then guaranteed to be able to choose a value for X_2 because the graph is **2-consistent**, for X_3 because it is **3-consistent**, and so on.
 - For each variable X_i , we need only **search through the d values** in the domain to find a value consistent with X_1, \dots, X_{i-1} .
- The total run time is only $O(n^2d)$.

Global Constraints

- A **global constraint** is one **involving** an **arbitrary** number of **variables** (but not necessarily all variables)
 - *Alldiff* constraint says that **all** the **variables** involved must **have distinct values**
- One simple form of **inconsistency** detection for *Alldiff* constraints works as follows:
 - if ***m* variables** are involved in the constraint, and if they have ***n*** possible **distinct values** altogether, and **$m > n$** , then the constraint **cannot be satisfied**.

Global Constraints

- Simple algorithm
 - Remove any variable in the constraint that has a singleton domain, and delete that variable's value from the domains of the remaining variables.
 - Repeat as long as there are singleton variables.
 - If at any point an empty domain is produced or there are more variables than domain values left, then an inconsistency has been detected.
- Can detect the inconsistency in the assignment $\{WA=\text{red}, NSW=\text{red}\}$
 - The variables SA , NT , and Q are connected by an *Alldiff* constraint
 - After applying AC-3 with the partial assignment, the domains of SA , NT , and Q are all reduced to {green,blue}.
 - We have three variables and only two colors, so the *Alldiff* constraint is violated.
- A simple consistency procedure for a higher-order constraint is sometimes more effective than applying arc-consistency to an equivalent set of binary constraints.



Sudoku

- A **Sudoku** board consists of **81 squares**, some of which are initially filled with **digits** from **1 to 9**.
- The puzzle is to **fill** in all the **remaining squares** such that **no digit appears twice in any row, column, or 3×3 box**.
- A row, column, or box is called a **unit**.

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| A | | | 3 | | 2 | | 6 | | |
| B | 9 | | | 3 | | 5 | | | 1 |
| C | | | 1 | 8 | | 6 | 4 | | |
| D | | | 8 | 1 | | 2 | 9 | | |
| E | 7 | | | | | | | | 8 |
| F | | | 6 | 7 | | 8 | 2 | | |
| G | | | 2 | 6 | | 9 | 5 | | |
| H | 8 | | | 2 | | 3 | | | 9 |
| I | | | 5 | | 1 | | 3 | | |

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| A | 4 | 8 | 3 | 9 | 2 | 1 | 6 | 5 | 7 |
| B | 9 | 6 | 7 | 3 | 4 | 5 | 8 | 2 | 1 |
| C | 2 | 5 | 1 | 8 | 7 | 6 | 4 | 9 | 3 |
| D | 5 | 4 | 8 | 1 | 3 | 2 | 9 | 7 | 6 |
| E | 7 | 2 | 9 | 5 | 6 | 4 | 1 | 3 | 8 |
| F | 1 | 3 | 6 | 7 | 9 | 8 | 2 | 4 | 5 |
| G | 3 | 7 | 2 | 6 | 8 | 9 | 5 | 1 | 4 |
| H | 8 | 1 | 4 | 2 | 5 | 3 | 7 | 6 | 9 |
| I | 6 | 9 | 5 | 4 | 1 | 7 | 3 | 8 | 2 |

Sudoku

- A **Sudoku** puzzle can be considered a **CSP** with **81 variables**, one for each square.
- We use the **variable names** **A1** through **A9** for the **top row** (left to right), down to **I1** through **I9** for the **bottom row**.
- The empty squares have the **domain** **{1,2,3,4,5,6,7,8,9}** and the **pre-filled squares** have a **domain** consisting of **a single value**.
- There are **27** different ***Alldiff*** constraints, one for each unit (row, column, and box of 9 squares):

$Alldiff(A1, A2, A3, A4, A5, A6, A7, A8, A9)$

$Alldiff(B1, B2, B3, B4, B5, B6, B7, B8, B9)$

...

$Alldiff(A1, B1, C1, D1, E1, F1, G1, H1, I1)$

$Alldiff(A2, B2, C2, D2, E2, F2, G2, H2, I2)$

...

$Alldiff(A1, A2, A3, B1, B2, B3, C1, C2, C3)$

$Alldiff(A4, A5, A6, B4, B5, B6, C4, C5, C6)$

...

Sudoku

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| A | | 3 | 2 | | 6 | | | | |
| B | 9 | | 3 | | 5 | | | | 1 |
| C | | 1 | 8 | | 6 | 4 | | | |
| D | | 8 | 1 | | 2 | 9 | | | |
| E | 7 | | | | | | | 8 | |
| F | | 6 | 7 | | 8 | 2 | | | |
| G | | 2 | 6 | | 9 | 5 | | | |
| H | 8 | | 2 | | 3 | | | | 9 |
| I | | 5 | 1 | | 3 | | | | |

| | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|---|---|---|---|---|---|---|---|---|---|
| A | 4 | 8 | 3 | 9 | 2 | 1 | 6 | 5 | 7 |
| B | 9 | 6 | 7 | 3 | 4 | 5 | 8 | 2 | 1 |
| C | 2 | 5 | 1 | 8 | 7 | 6 | 4 | 9 | 3 |
| D | 5 | 4 | 8 | 1 | 3 | 2 | 9 | 7 | 6 |
| E | 7 | 2 | 9 | 5 | 6 | 4 | 1 | 3 | 8 |
| F | 1 | 3 | 6 | 7 | 9 | 8 | 2 | 4 | 5 |
| G | 3 | 7 | 2 | 6 | 8 | 9 | 5 | 1 | 4 |
| H | 8 | 1 | 4 | 2 | 5 | 3 | 7 | 6 | 9 |
| I | 6 | 9 | 5 | 4 | 1 | 7 | 3 | 8 | 2 |

- *Alldiff constraints* have been expanded into **binary constraints** (such as $A1 \neq A2$) so that we can apply the **AC-3** algorithm directly.
 - Consider variable **E6**,
 - from the constraints **in the box**, we can **remove** 1, 2, 7, and 8 from E6's domain.
 - from the constraints **in its column**, we can **eliminate** 5, 6, 2, 8, 9, and 3.
 - that leaves **E6** with a domain of {4}
 - Consider variable **I6**
 - applying arc consistency **in its column**, we eliminate 5, 6, 2, 4 (E6 must be 4), 8, 9, and 3.
 - we eliminate 1 by arc consistency with I5, and we are left with only the value 7 in the domain of I6.
 - now there are 8 known values in column 6, so arc consistency can infer that A6 must be 1.
- Inference continues along these lines, and eventually, AC-3 can solve the entire puzzle—all the **variables** have their domains **reduced to a single value**

Backtracking Search for CSPs

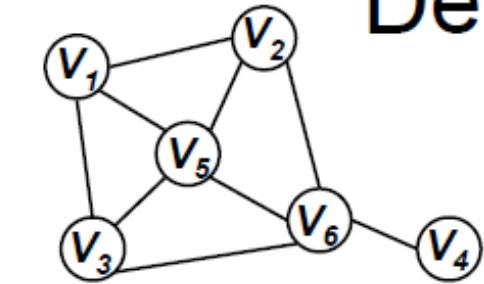
- Sometimes we can **finish** the **constraint propagation** process and still have **variables** with **multiple possible values**.
- In that case we have to **search** for a **solution**.

Backtracking Search for CSPs



- Consider **how** a standard **depth-limited search** could **solve CSPs**.
- A state would be a partial assignment, and an action would extend the assignment, **$NSW = red$** or **$SA = blue$**
- For a CSP with **n** variables of domain size **d**
 - We have a **search tree** where all the complete **assignments** are **leaf nodes** at **depth n** .
 - But notice that the **branching factor** at the **top level** would be **nd** because any of **d** values can be assigned to any of **n variables**.
 - At the next level, the branching factor is **$(n - 1)d$** , and so on for **n levels**.
 - So the **tree has $n! \cdot d^n$ leaves**, even though there are only **d^n** possible complete assignments!

Depth First Search



| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| ? | ? | ? | ? | ? | ? |

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | ? | ? | ? | ? | ? |

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| R | ? | ? | ? | ? | ? |

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| G | ? | ? | ? | ? | ? |

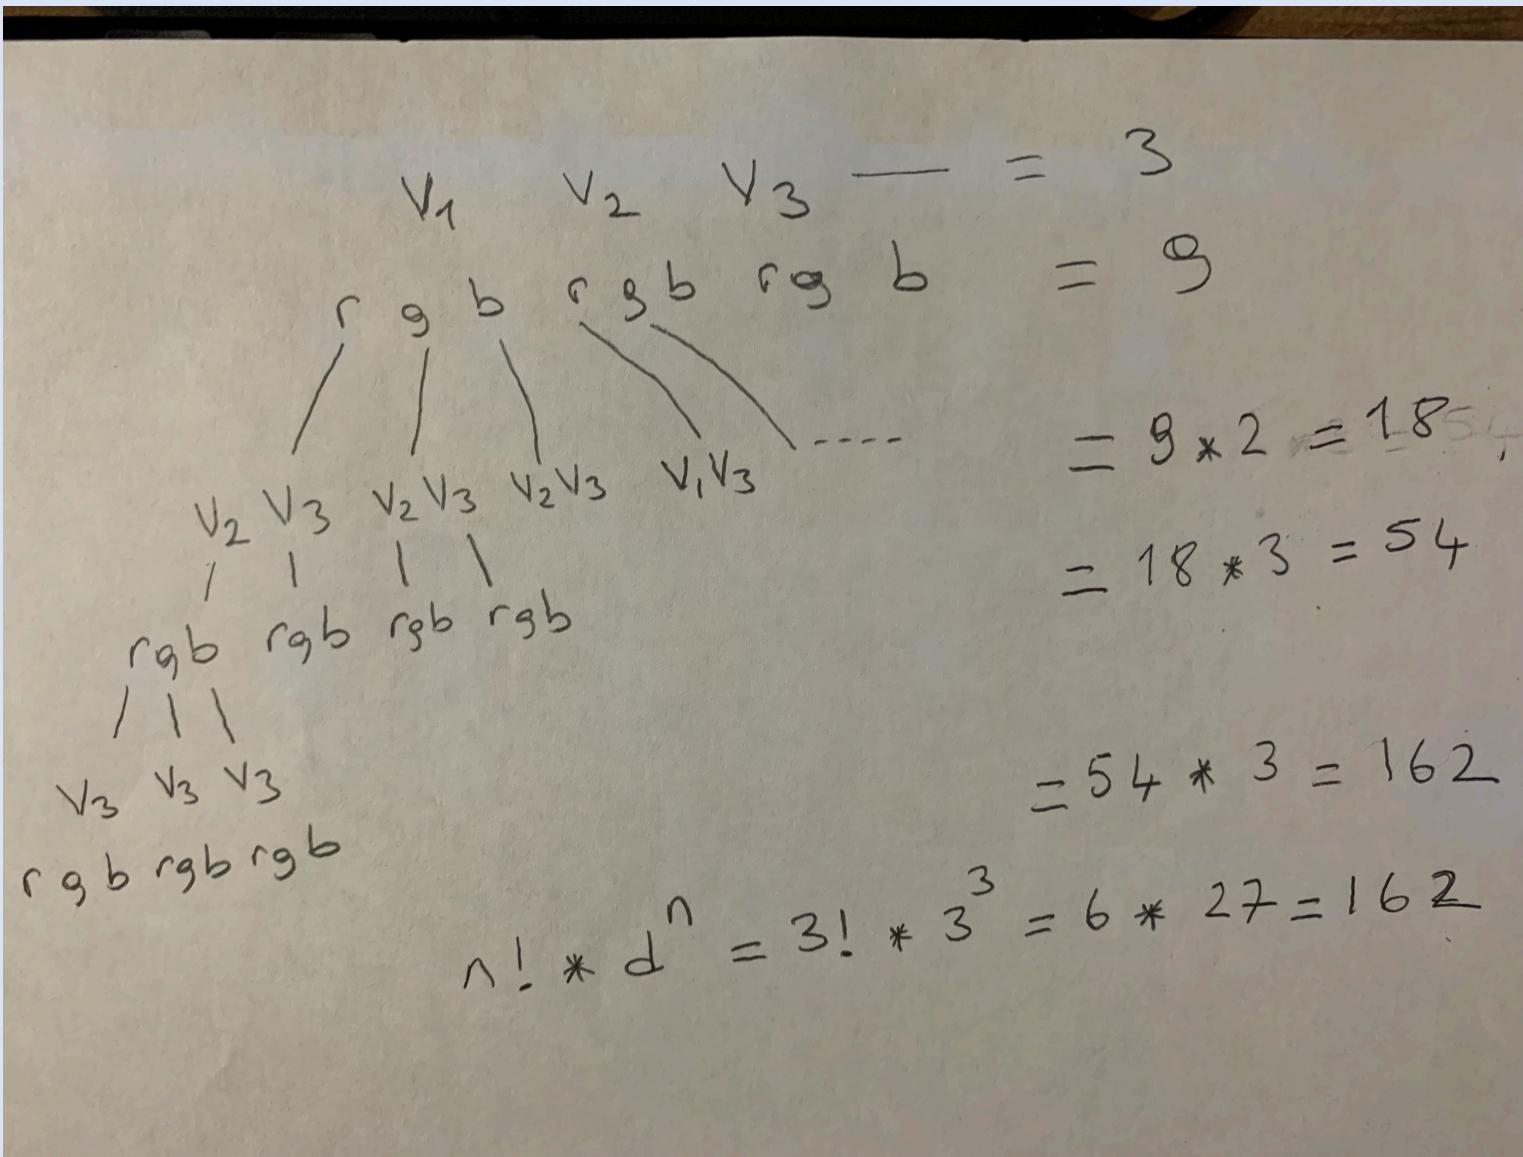
| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | B | ? | ? | ? | ? |

• Recursively:

- For every possible value in D :
 - Set the next unassigned variable in the successor to that value
 - Evaluate the successor of the current state with this variable assignment
 - Stop as soon as a solution is found

9d

Search tree for CSP

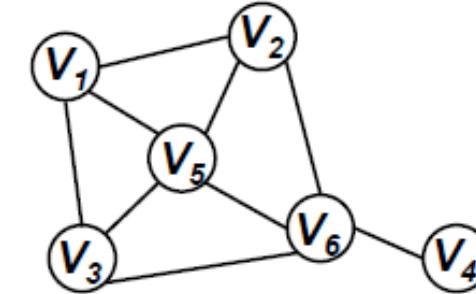


Backtracking Search for CSPs

- We can get back that factor of $n!$ by recognizing a crucial property of CSPs: *commutativity*.
- In CSPs, it makes no difference if we first assign NSW = red and then SA = blue, or the other way around.
- Only consider a single variable at each node in the search tree.
- At the root we might make a choice between SA=red, SA=green, and SA=blue, but we would never choose between NSW=red and SA=blue.
- With this restriction, the number of leaves is d^n

Backtracking DFS

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| ? | ? | ? | ? | ? | ? |



| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | ? | ? | ? | ? | ? |

Order of values:
 (B, R, G)

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | B | ? | ? | ? | ? |

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | R | ? | ? | ? | ? |

Don't even consider
that branch because
 $V_2 = B$ is inconsistent
with the parent state

| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | R | R | B | ? | ? |



| | | | | | |
|-------|-------|-------|-------|-------|-------|
| V_1 | V_2 | V_3 | V_4 | V_5 | V_6 |
| B | R | R | B | G | ? |

Backtrack to the
previous state
because no valid
assignment can
be found for V_6

Example – N Queen Problem

- <https://www.youtube.com/watch?v=0DeznFqrgAI>

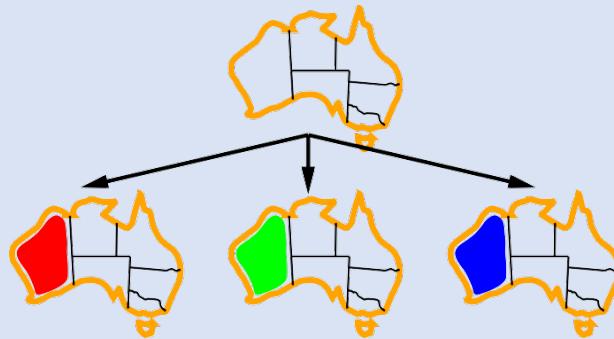
Backtracking Search for CSPs

- Depth-first search for CSPs with single-variable assignments is called backtracking search
- Backtracking search is the basic uninformed algorithm for CSPs

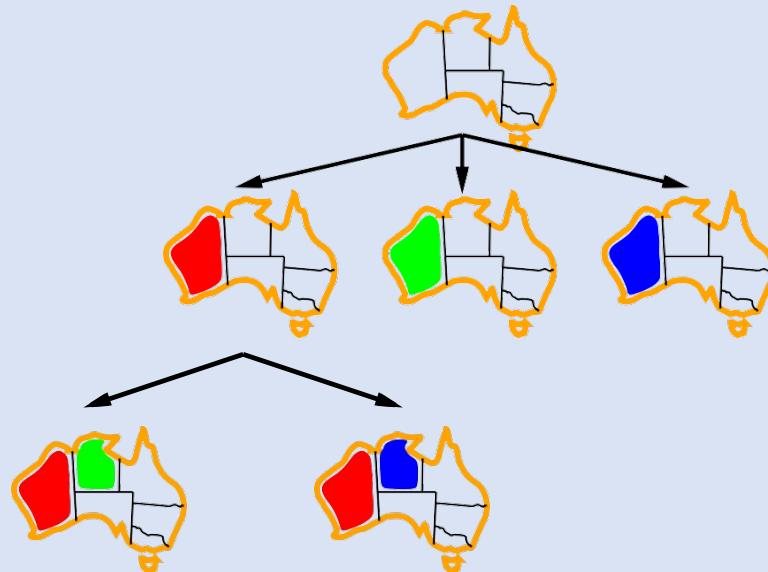
Backtracking example



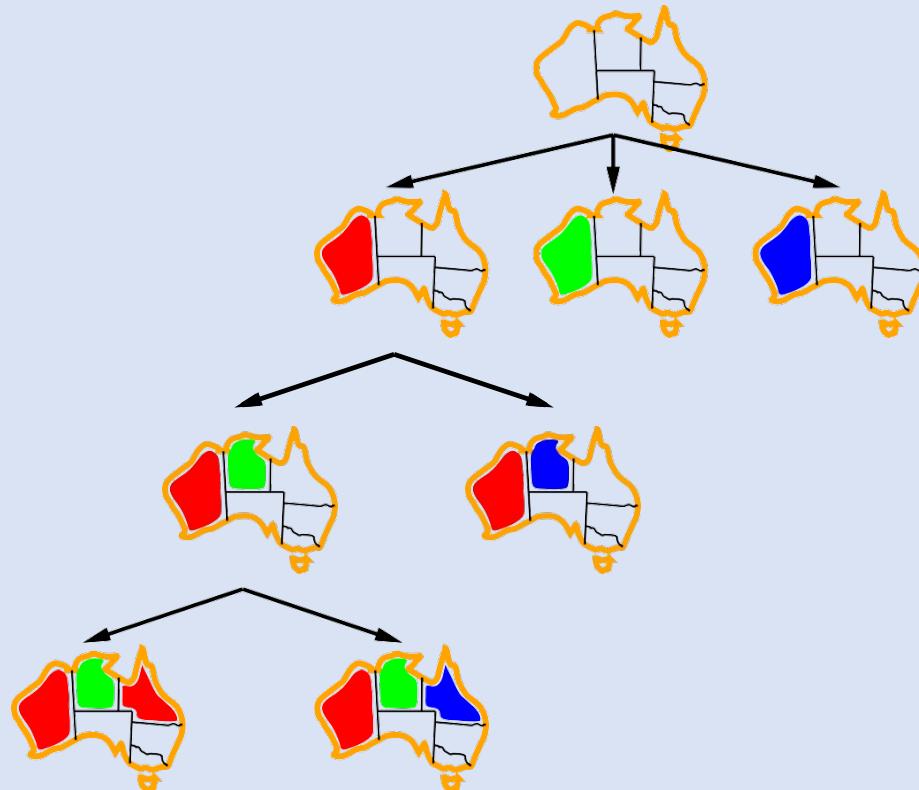
Backtracking example



Backtracking example



Backtracking example



Backtracking algorithm

```
function BACKTRACKING-SEARCH(csp) returns a solution or failure
    return BACKTRACK(csp, {})

function BACKTRACK(csp, assignment) returns a solution or failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp, assignment)
    for each value in ORDER-DOMAIN-VALUES(csp, var, assignment) do
        if value is consistent with assignment then
            add  $\{var = value\}$  to assignment
            inferences  $\leftarrow$  INFERENCE(csp, var, assignment)
            if inferences  $\neq$  failure then
                add inferences to csp
                result  $\leftarrow$  BACKTRACK(csp, assignment)
                if result  $\neq$  failure then return result
                remove inferences from csp
                remove  $\{var = value\}$  from assignment
    return failure
```

Figure 5.5 A simple backtracking algorithm for constraint satisfaction problems. The algorithm is modeled on the recursive depth-first search of Chapter 3. The functions SELECT-UNASSIGNED-VARIABLE and ORDER-DOMAIN-VALUES implement the general-purpose heuristics discussed in Section 5.3.1. The INFERENCE function can optionally impose arc-, path-, or k -consistency, as desired. If a value choice leads to failure (noticed either by INFERENCE or by BACKTRACK), then value assignments (including those made by INFERENCE) are retracted and a new value is tried.

Variable and value ordering

- The backtracking algorithm contains the line

$var \leftarrow \text{SELECT-UNASSIGNED-VARIABLE}(csp, assignment)$.

- Strategies for SELECT-UNASSIGNED-VARIABLE

- static ordering: choose the variables in order, $\{X_1, X_2, \dots\}$.
- choose randomly.

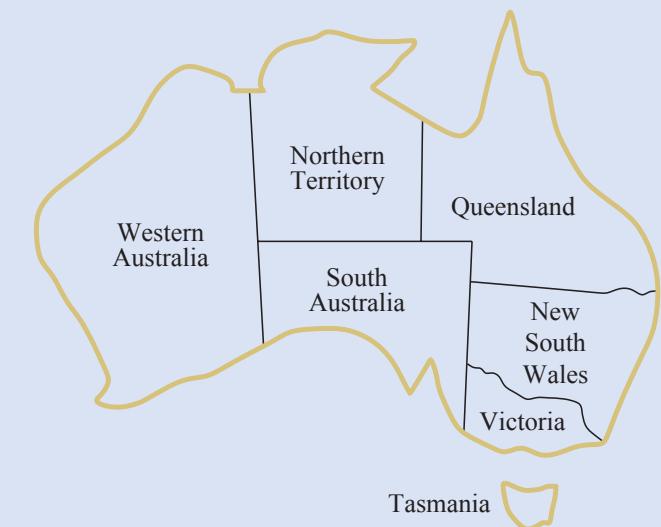
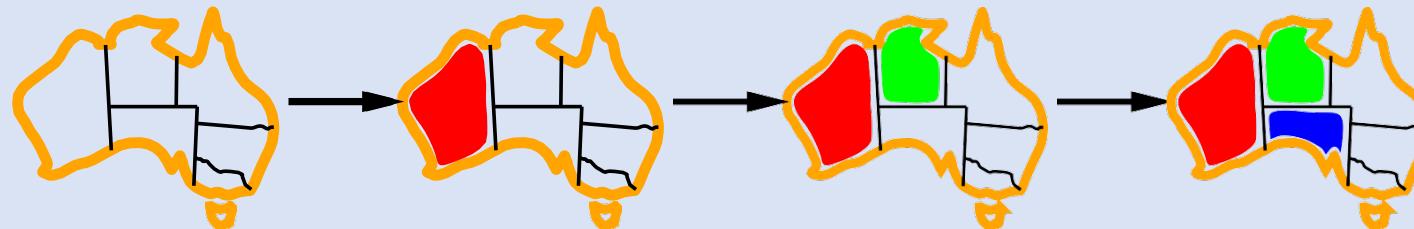
- Neither strategy is optimal.

- after the assignments for $WA=\text{red}$ and $NT=\text{green}$, there is only one possible value for SA , so it makes sense to assign $SA=\text{blue}$ next rather than assigning Q .
- after SA is assigned, the choices for Q , NSW , and V are all forced.



Minimum remaining values

Minimum remaining values (MRV):
choose the variable with the fewest legal values



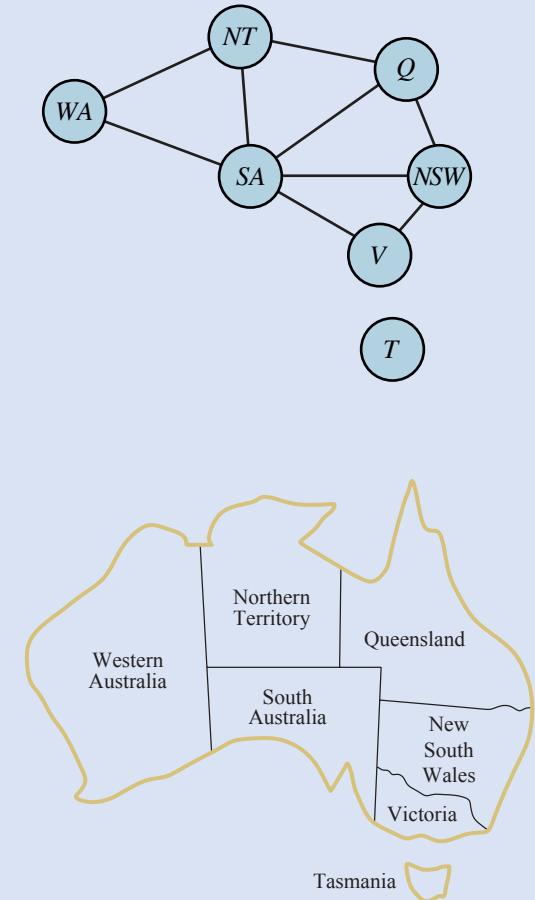
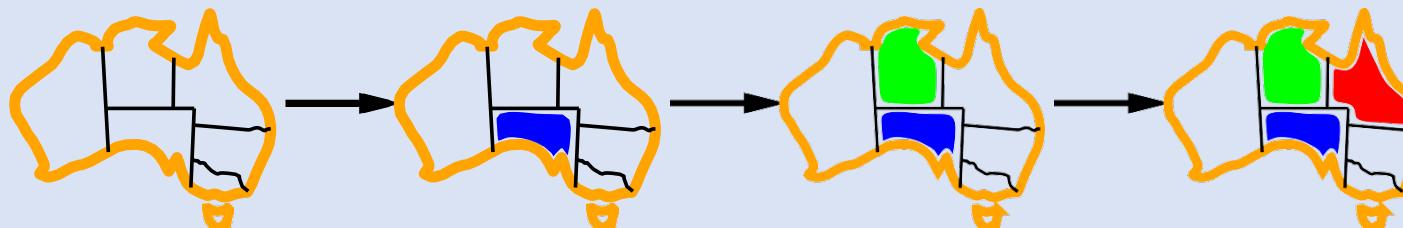
Degree heuristic

- The **MRV** heuristic **doesn't help** at all in choosing the **first region** to color in Australia, because initially **every region** has **three legal colors**.

Tie-breaker among MRV variables

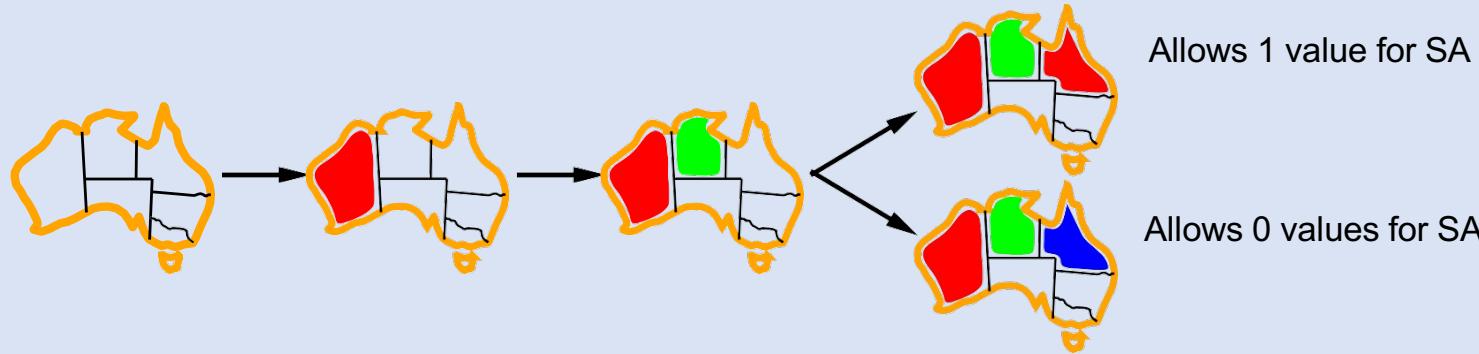
Degree heuristic:

choose the variable with the most constraints on remaining variables



Least constraining value

Given a variable, choose the least constraining value:
the one that rules out the fewest values in the remaining variables



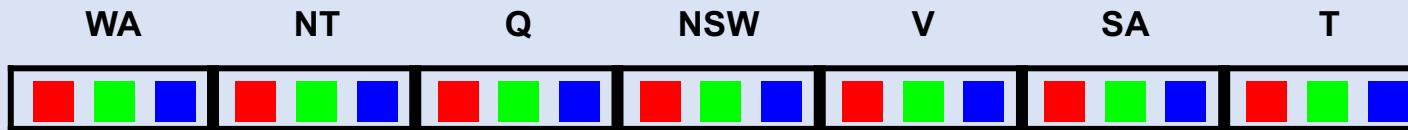
Combining these heuristics makes 1000 queens feasible



INFERENCE - Forward checking

Idea: Keep track of remaining legal values for unassigned variables (arc consistency)

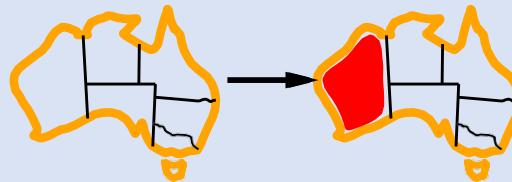
Terminate search when any variable has no legal values



INFERENCE - Forward checking

Idea: Keep track of remaining legal values for unassigned variables (arc consistency)

Terminate search when any variable has no legal values



| WA | NT | Q | NSW | V | SA | T |
|----------------------|----------------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| Red
Green
Blue |
| Red | | Green
Blue | Red
Green
Blue | Red
Green
Blue | Green
Blue | Red
Green
Blue |



INFERENCE - Forward checking

Idea: Keep track of remaining legal values for unassigned variables (arc consistency)

Terminate search when any variable has no legal values



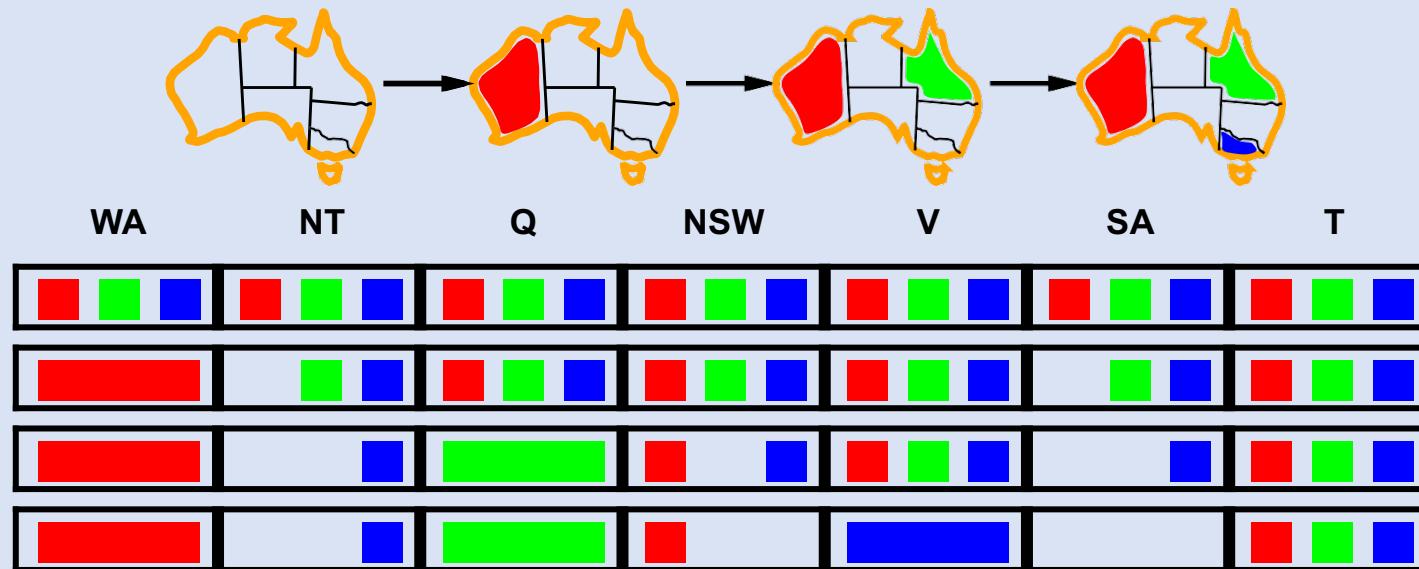
| WA | NT | Q | NSW | V | SA | T |
|----------------------|----------------------|----------------------|----------------------|----------------------|----------------------|----------------------|
| Red
Green
Blue |
| Red | | Green
Blue | Red
Green
Blue | Red
Green
Blue | Green
Blue | Red
Green
Blue |
| Red | | Blue | Green | Red
Green
Blue | Blue | Red
Green
Blue |



INFERENCE - Forward checking

Idea: Keep track of remaining legal values for unassigned variables (arc consistency)

Terminate search when any variable has no legal values



- MAC detects inconsistencies like this. After a variable X_i is assigned a value, the INFERENCE procedure calls AC-3, but instead of a queue of all arcs in the CSP, we start with only the arcs (X_j, X_i) for all X_j that are unassigned variables that are neighbors of X_i .
- From there, AC-3 does constraint propagation in the usual way, and if any variable has its domain reduced to the empty set, the call to AC-3 fails and we know to backtrack immediately.

Intelligent backtracking: Looking backward



- Chronological backtracking :
 - what to do when a branch of the search fails: back up to the preceding variable and try a different value for it.
 - the most recent decision point is revisited.
- Consider a fixed variable ordering Q, NSW, V, T, SA, WA, NT
 - partial assignment {Q=red, NSW=green, V=blue, T=red}
 - when we try the next variable, SA , we see that every value violates a constraint.
 - we back up to T and try a new color for Tasmania!
 - Obviously this is silly - recoloring Tasmania cannot possibly help in resolving the problem with South Australia.

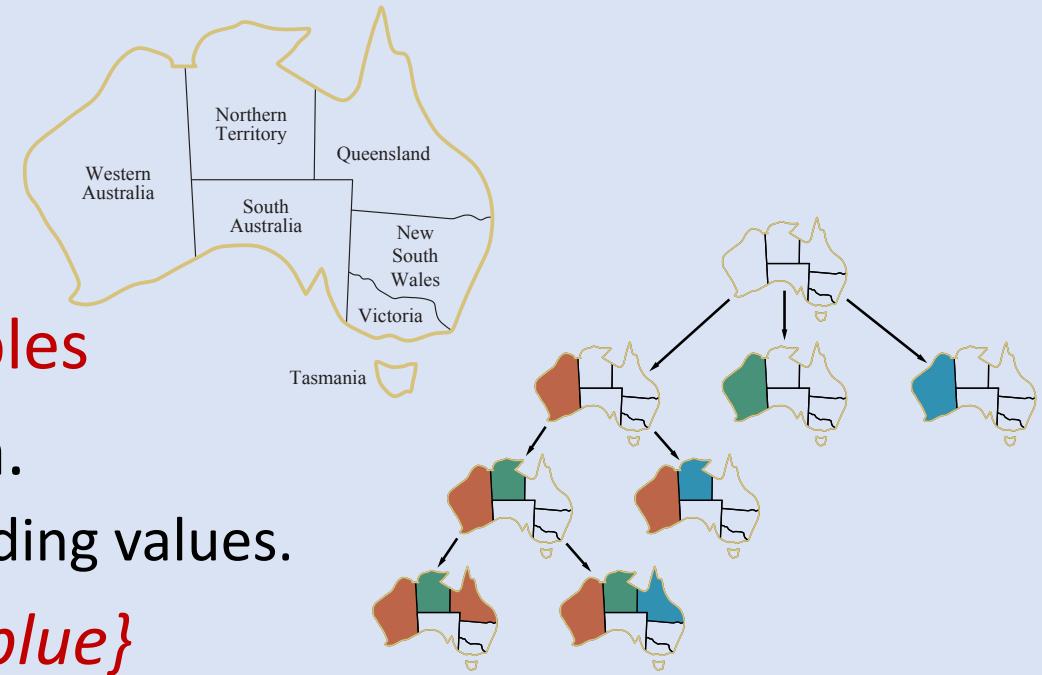
Intelligent backtracking: Looking backward



- Backtrack to a variable that might fix the problem
 - a variable that was responsible for making one of the possible values of SA impossible.
 - keep track of a set of assignments that are in conflict with some value for SA.
 - The conflict set: $\{Q=\text{red}, NSW=\text{green}, V=\text{blue}\}$.
 - the backjumping method backtracks to the most recent assignment in the conflict set;
 - backjumping would jump over Tasmania and try a new value for $V(Q, NSW, V, T, SA, WA, NT)$.
 - It is easily implemented by a modification to BACKTRACK such that it accumulates the conflict set while checking for a legal value to assign.
 - If no legal value is found, the algorithm should return the most recent element of the conflict set along with the failure indicator.

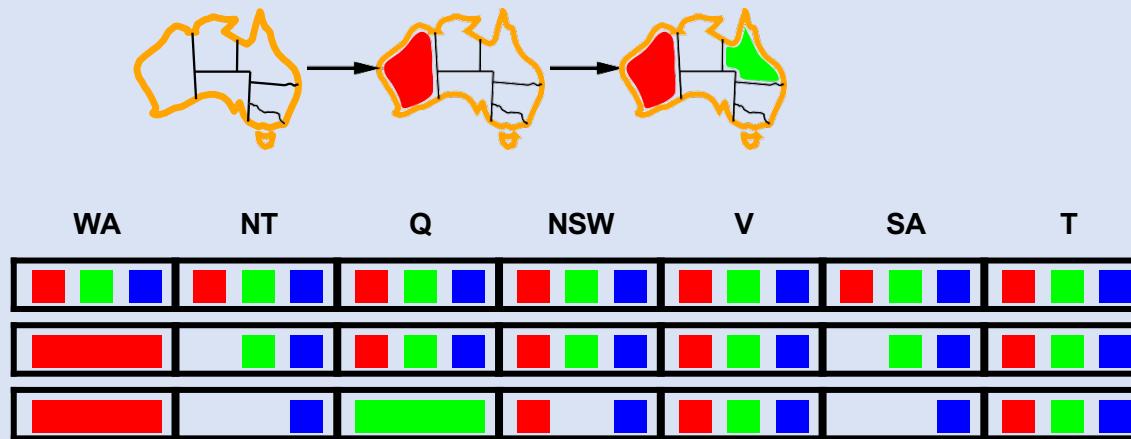
Constraint learning

- The idea of **finding a minimum set of variables** from the **conflict set** that causes the problem.
 - **no-good**: set of variables, with their corresponding values.
- Consider the state $\{WA=\text{red}, NT=\text{green}, Q=\text{blue}\}$
- **Forward checking** can tell us this state is a **no-good** because there is **no valid assignment to SA**.
 - suppose that the search tree were actually part of a larger search tree that started by first assigning values for V and T .
 - it would be worthwhile to record $\{WA=\text{red}, NT=\text{green}, Q=\text{blue}\}$ as a **no-good** because we are going to run into the **same problem** again for **each possible** set of assignments to V and T .



Constraint propagation

Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:



NT and *SA* cannot both be blue!

Constraint propagation repeatedly enforces constraints locally

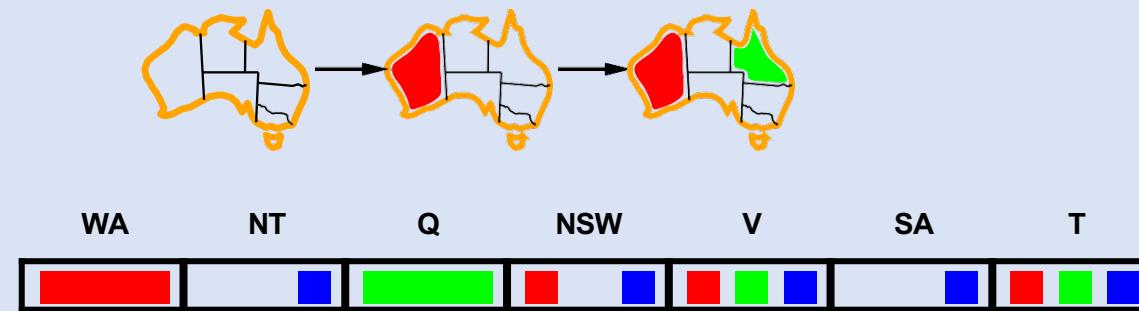


Arc consistency

Simplest form of propagation makes each arc **consistent**

$X \rightarrow Y$ is consistent iff

for every value x of X there is some allowed y

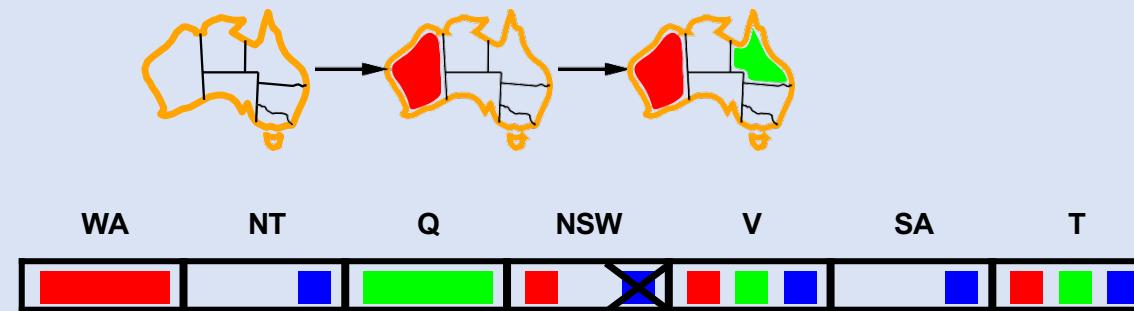


Arc consistency

Simplest form of propagation makes each arc **consistent**

$X \rightarrow Y$ is consistent iff

for every value x of X there is some allowed y

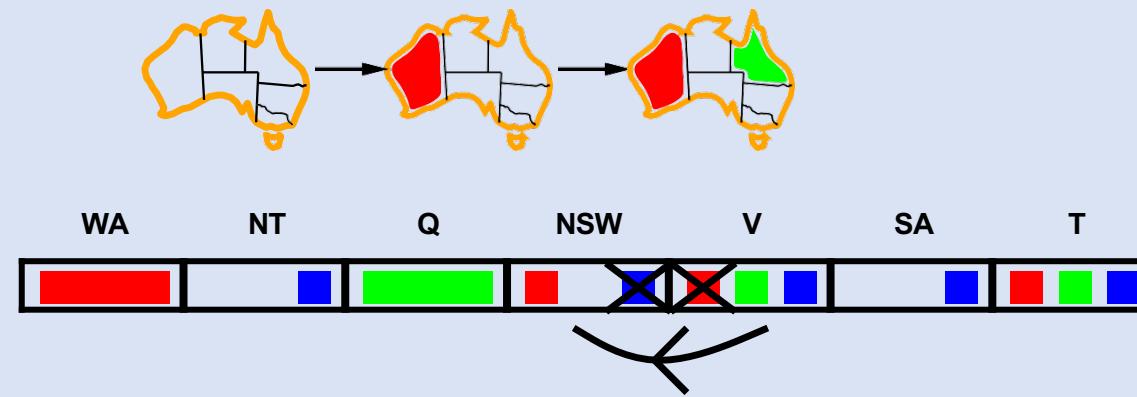


Arc consistency

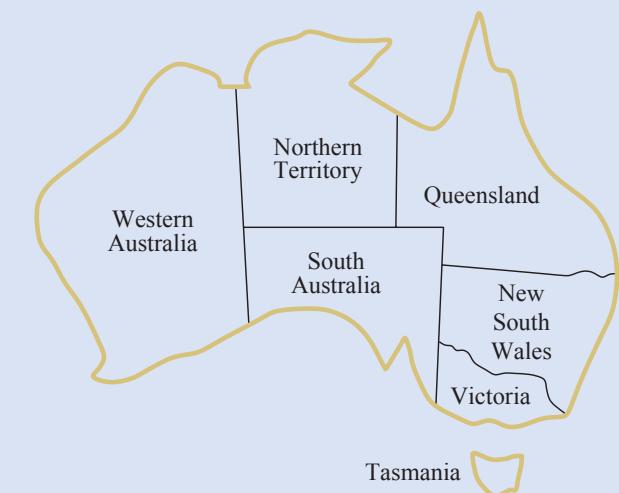
Simplest form of propagation makes each arc **consistent**

$X \rightarrow Y$ is consistent iff

for every value x of X there is some allowed y



If X loses a value, neighbors of X need to be rechecked

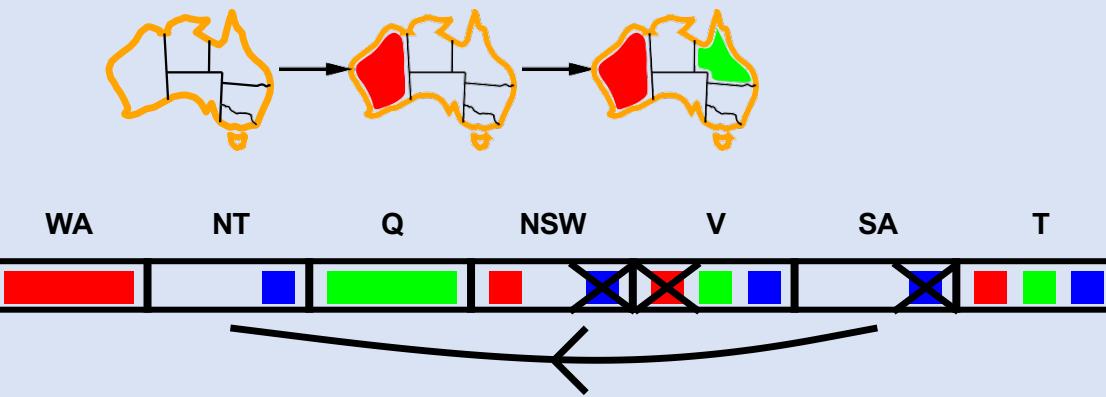


Arc consistency

Simplest form of propagation makes each arc **consistent**

$X \rightarrow Y$ is consistent iff

for **every** value x of X there is **some** allowed y



If X loses a value, neighbors of X need to be rechecked

Arc consistency detects failure earlier than forward checking

Can be run as a preprocessor or after each assignment



Arc consistency algorithm

```
function AC-3(csp) returns the CSP, possibly with reduced domains
    inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
    local variables: queue, a queue of arcs, initially all the arcs in csp

    while queue is not empty do
         $(X_i, X_j) \leftarrow \text{Remove-First}(\text{queue})$ 
        if Remove-Inconsistent-Values( $X_i, X_j$ ) then
            for each  $X_k$  in Neighbors[ $X_i$ ] do
                add  $(X_k, X_i)$  to queue



---


function Remove-Inconsistent-Values( $X_i, X_j$ ) returns true iff succeeds
    removed  $\leftarrow \text{false}$ 
    for each  $x$  in Domain[ $X_i$ ] do
        if no value  $y$  in Domain[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
            then delete  $x$  from Domain[ $X_i$ ]; removed  $\leftarrow \text{true}$ 
    return removed
```

$O(n^2d^3)$, can be reduced to $O(n^2d^2)$ (but detecting **all** is NP-hard)

The End!

