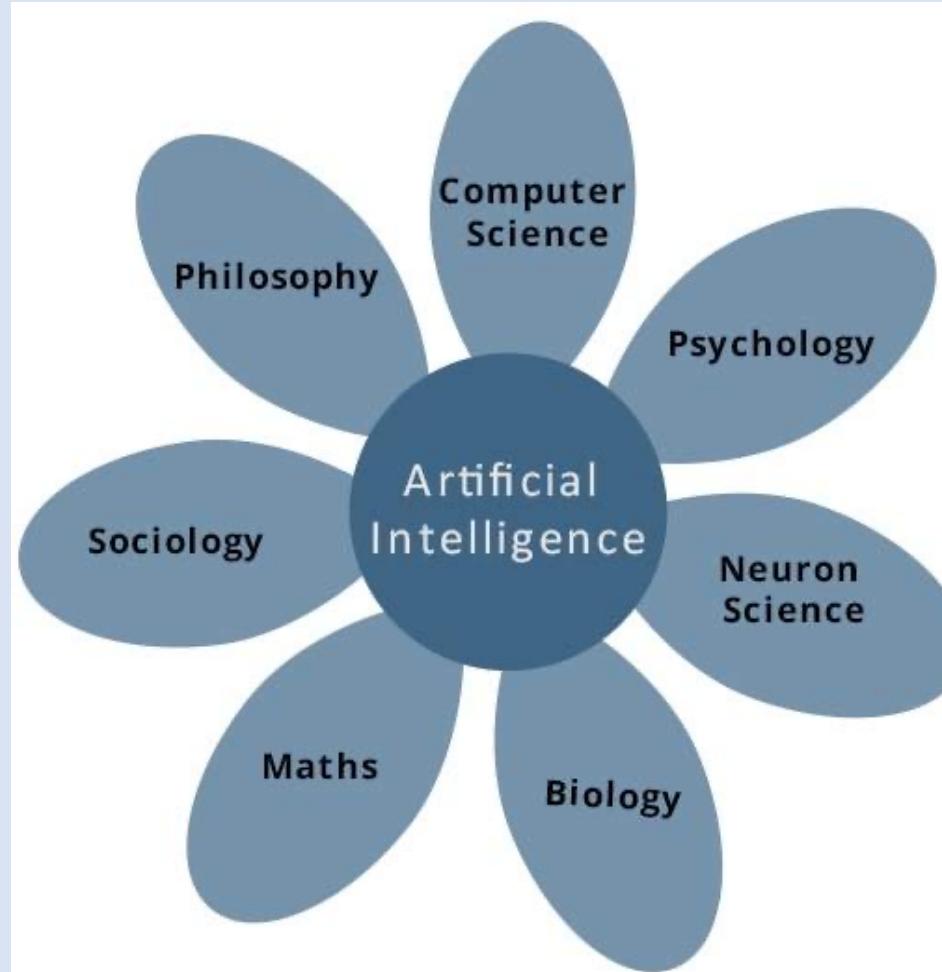


Artificial Intelligence

Multi-Agent Systems (MAS)

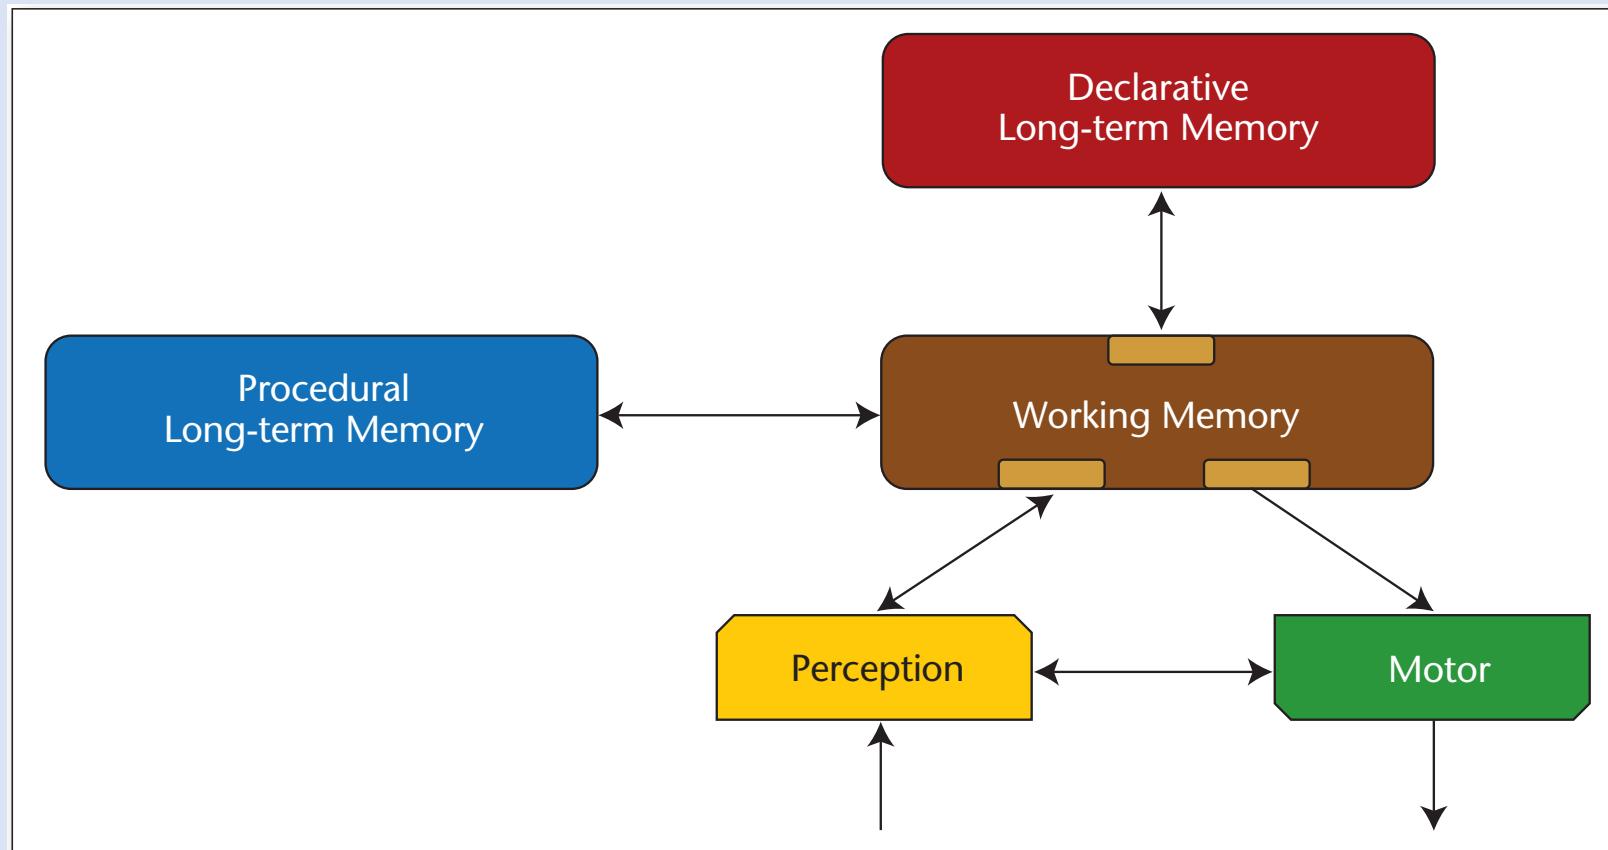
Dr. Bilgin Avenoğlu

AI Disciplines



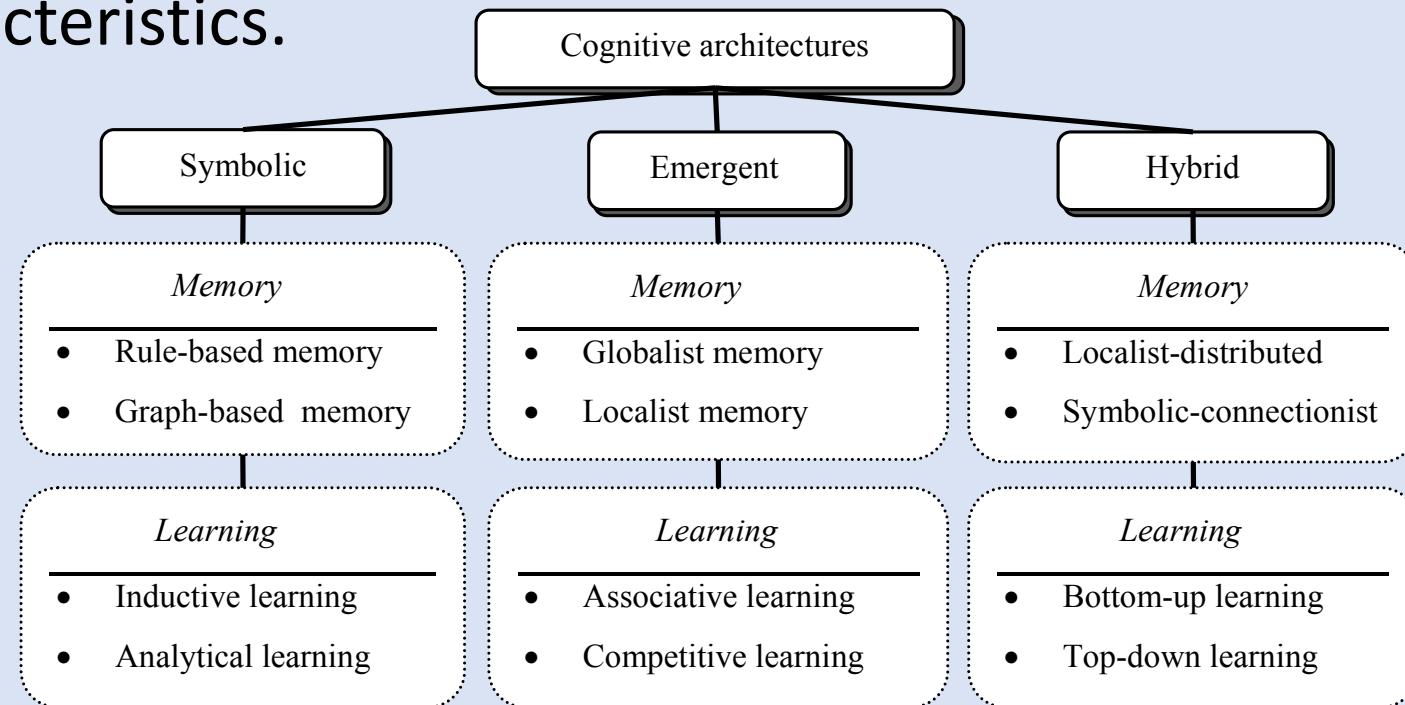
A Standard Model of the Mind

- The mind is not an undifferentiated pool of information and processing, but is built of **independent modules** that have **distinct functionalities**.



Cognitive Architectures

- Cognitive architecture (CA) research is a theme in AI that models the main factors participated in **our thinking** and **decision** and concentrates on the **relationships** among them.
- CA mostly refers to the **computational model** simulating **human's cognitive** and **behavioral** characteristics.



Abbreviation	Planning	Motivation	Interaction	Proposed year				
Perception	Attention	Emotion						
Memory	Learning							
Actuation	Reasoning							
Perception	Memory	Attention	Emotion	Interaction				
4CAPS	✓			2006				
4D/RCS	✓			2000				
ACT-R	✓	✓	✓	2004				
ADAPT	✓	✓		2004				
ARCADIA	✓	✓	✓	2015				
ARDIS	✓	✓		2009				
ASMO	✓	✓	✓	2010				
BBD	✓	✓		2005				
BECCA	✓	✓	✓	2009				
BDI	✓	✓	✓	1998				
CARACaS	✓	✓		✓	2011			
Casimir	✓		✓	✓	2011			
CELTS	✓	✓	✓	✓	2011			
CERA-CRANIUM	✓	✓	✓	✓	✓	2009		
Cerebus	✓	✓	✓		✓	2000		
CHARISMA	✓	✓	✓		✓	✓	2011	
CHREST	✓		✓	✓			1992	
CLARION	✓	✓	✓	✓			1998	
COGNET/iGEN	✓	✓	✓	✓	✓		2000	
Cognitive Symmetry Engine	✓	✓	✓				2011	
CoJACK	✓	✓	✓	✓	✓	✓	2009	
Companions	✓		✓	✓		✓	2006	
Cosy	✓	✓	✓	✓		✓	✓	2004
CORTEX	✓	✓	✓	✓		✓		2014
DAC	✓	✓	✓	✓	✓		✓	2012
DIARC	✓	✓	✓	✓	✓	✓		2006
DiPRA	✓	✓	✓	✓		✓		2005
DSO-CA	✓	✓	✓	✓				2011
EMILE	✓	✓	✓	✓	✓	✓		2000
EPAM	✓	✓	✓	✓				1995
EPIC	✓	✓	✓	✓				1994

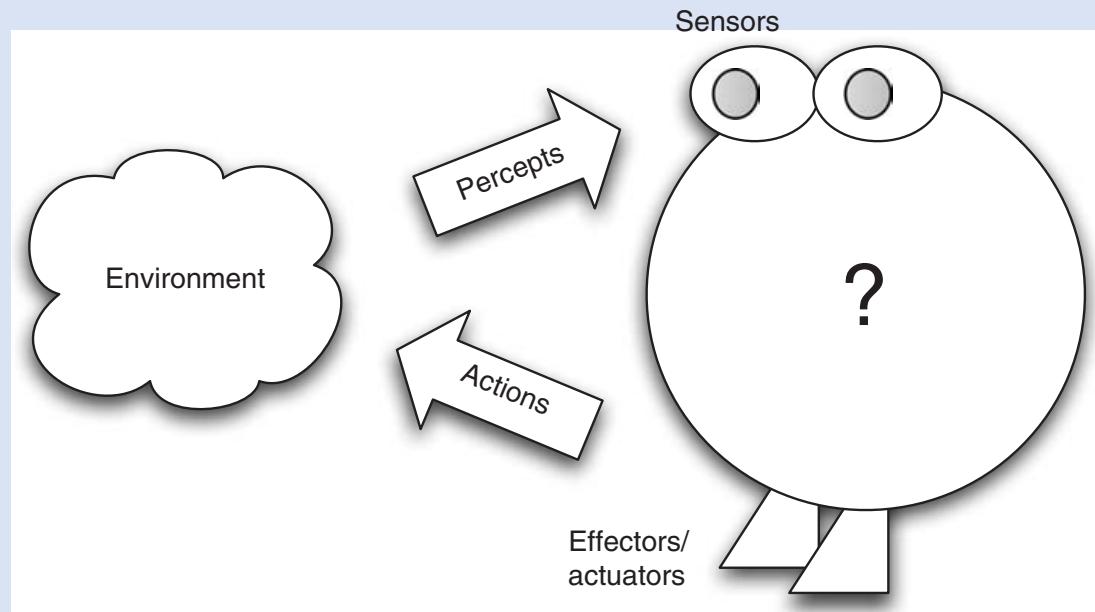
GLAIR	✓	✓		✓			✓		2009	
GMU-BICA	✓		✓	✓					2005	
HTM	✓	✓	✓		✓				2007	
ICARUS	✓	✓	✓	✓					2004	
iCub	✓	✓	✓			✓	✓		2004	
Ikon Flux	✓		✓	✓	✓				2009	
IMPRINT	✓		✓	✓					2002	
Leabra	✓		✓						2012	
LIDA	✓	✓	✓			✓			2006	
MAMID	✓	✓	✓	✓		✓	✓		2000	
MDB	✓	✓	✓		✓				2006	
MicroPsi	✓	✓	✓		✓		✓		✓	2007
MusiCog	✓		✓		✓				2012	
NARS			✓	✓	✓				2013	
Novamente Engine	✓	✓	✓	✓					2007	
OpenCogPrime	✓	✓	✓	✓	✓	✓	✓	✓	✓	2013
PMFserv			✓			✓		✓		2001
Pogamut	✓	✓	✓							2009
Polyscheme	✓		✓	✓		✓				2005
R-CAST	✓	✓	✓	✓				✓		2008
REAPER	✓	✓	✓					✓		2001
REM	✓	✓	✓	✓						2001
SAL	✓	✓	✓	✓						2008
SEMLC	✓	✓	✓		✓			✓		2011
Shruti	✓	✓	✓	✓						2000
SiMA	✓	✓	✓	✓					✓	2013
Soar			✓	✓	✓					1987

MAS Definition

- Types of software:
 - functional programs: takes some **input**, **chews** over this input, and then on the basis of this, produces some **output** and **halts**.
 - ‘read a list of numbers and print the average’
- Unfortunately, many programs **do not have** this simple *input-compute-output* operational structure.
 - many of the systems we need to build in practice have a ‘**reactive**’ flavour, in the sense that they have to **maintain a long- term, ongoing interaction** with their **environment**;
 - operating systems, process control systems, online banking systems, web servers
- Agent-based systems: **active**, **purposeful** **producers of actions**
 - they are **sent out into their environment** to **achieve goals** for us, and we want them to **actively pursue these goals**, figuring out for themselves **how best to accomplish these goals**, rather than having to be told in **low-level detail** how to do it.

MAS Definition

- Agent systems are *situated* in some *environment*.
- Agents are capable of *sensing* their environment (*via sensors*),
- Agents have a repertoire of possible *actions* that they can perform (*via effectors* or *actuators*) in order to modify their *environment*.



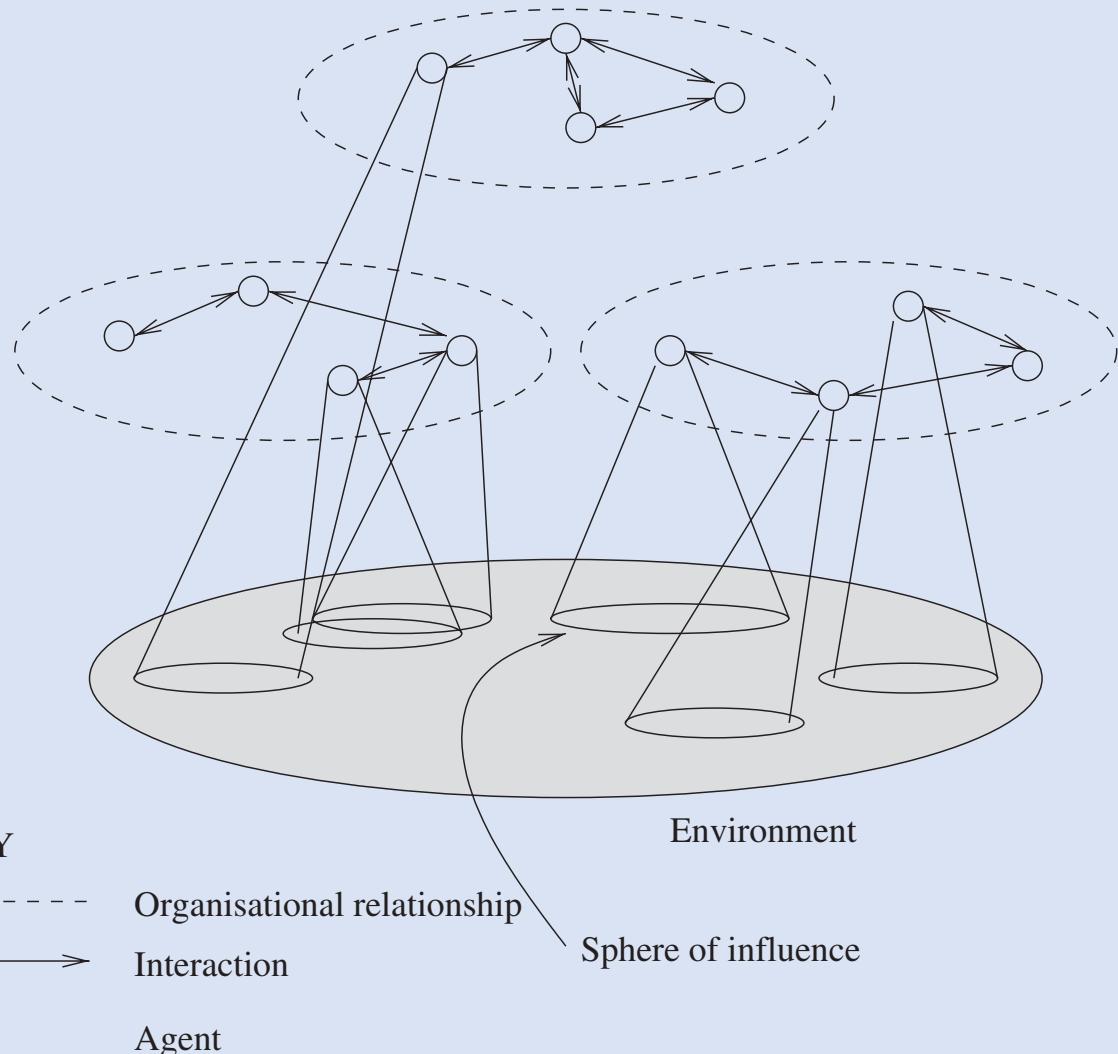
MAS Definition

- agents should have the following properties:
 - autonomy;
 - proactiveness;
 - generating and attempting to achieve goals; not driven solely by events; taking the initiative.
 - reactivity;
 - If a program's environment is guaranteed to be fixed, a program can just execute blindly.
 - social ability.
 - Social ability in agents is the ability to interact with other agents (and possibly humans) via **cooperation**, **coordination**, and **negotiation**.

Level	Definition
1	The computer offers no assistance, human must do it all
2	The computer offers a complete set of action alternatives, and
3	Narrows the selection down to a few, or
4	Suggests one, and
5	Executes that suggestion if the human approves, or
6	Allows the human a restricted time to veto before automatic execution, or
7	Executes automatically, then necessarily informs the human, or
8	Informs him after execution only if he asks, or
9	Informs him after execution if it, the computer, decides to
10	The computer decides everything and acts autonomously, ignoring the human

MAS

- We have talked about agents occupying an **environment in isolation**.
- In practice, '**single agent systems**' are **rare**.
- The more common case is for agents to **inhabit an environment which contains other agents**, giving a **multi-agent system**.



The Belief-Desire-Intention (BDI) Agent Model

- The idea that we can talk about computer programs as if they have a ‘mental state’.
 - **Beliefs** are information the agent has about the world.
 - **Desires** are all the possible states of affairs that the agent might like to accomplish.
 - **Intentions** are the states of affairs that the agent has decided to work towards.
 - Intentions may be goals that are delegated to the agent, or may result from considering options: we think of an agent looking at its options and choosing between them.
- We can imagine our agent **starting with some delegated goal**, and then **considering the possible options** that are compatible with this delegated goal; the **options that it chooses** are then **intentions**, which the agent is committed to.

Practical Reasoning

- How does an agent with **beliefs, desires** and **intentions** go from these **to its actions**?
 - a matter of **weighing conflicting considerations** for and against **competing** options, where the relevant considerations are provided by what the agent desires/values/cares about and what the agent believes.
- Human practical reasoning seems to consist of two distinct activities: **deliberation** (fixing upon states of affairs that we want to achieve, i.e. our intentions); and **means-ends reasoning** (deciding how to act so as to bring about our intentions).

A Computational Model of BDI Practical Reasoning

```
1.  $B \leftarrow B_0;$  /*  $B_0$  are initial beliefs */
2.  $I \leftarrow I_0;$  /*  $I_0$  are initial intentions */
3. while true do
4.   get next percept  $\rho$  via sensors;
5.    $B \leftarrow brf(B, \rho);$ 
6.    $D \leftarrow options(B, I);$ 
7.    $I \leftarrow filter(B, D, I);$ 
8.    $\pi \leftarrow plan(B, I, Ac);$  /*  $Ac$  is the set of actions */
9.   while not (empty( $\pi$ ) or succeeded( $I, B$ ) or impossible( $I, B$ )) do
10.     $\alpha \leftarrow$  first element of  $\pi;$ 
11.    execute( $\alpha$ );
12.     $\pi \leftarrow$  tail of  $\pi;$ 
13.    observe environment to get next percept  $\rho;$ 
14.     $B \leftarrow brf(B, \rho);$ 
15.    if reconsider( $I, B$ ) then
16.       $D \leftarrow options(B, I);$ 
17.       $I \leftarrow filter(B, D, I);$ 
18.    end-if
19.    if not sound( $\pi, I, B$ ) then
20.       $\pi \leftarrow plan(B, I, Ac)$ 
21.    end-if
22.  end-while
23. end-while
```

Figure 2.1 Overall control loop for a BDI practical reasoning agent.

The Procedural Reasoning System (PRS)

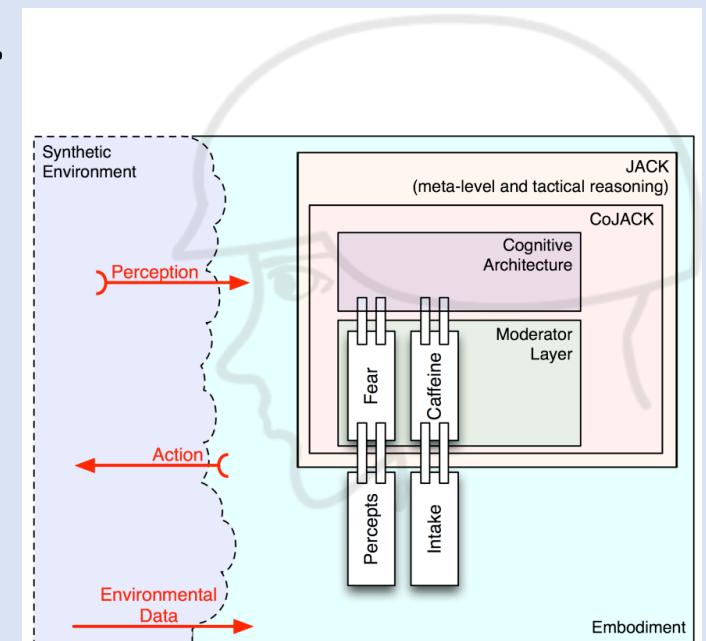
- PRS is originally developed at Stanford Research Institute by Michael Georgeff and Amy Lansky, was perhaps the first agent architecture to explicitly embody the BDI paradigm, and has proved to be one of the most durable approaches to developing agents to date.
- PRS has been re-implemented several times since the mid-1980s,
 - JACK is a commercially available programming language of PRS ,
 - Agent Oriented Software – the Australian AI company since 1997.



CoJACK

- CoJACK™ is cognitive architecture used for modelling the variation in human behaviour.
- It is used in simulation systems to underpin virtual actors.
- Humans share a common cognitive architecture and physiology, and that variation results from individual differences in knowledge and the values of the architecture's and physiology's parameters.

differences across individuals, differences across time on a task within an individual, and the effects of context on a task, such as time pressure, heat, sleep patterns, and food, drink, and psychoactive substances such as caffeine.



AgentSpeak and Jason

- In 1996, Anand Rao created a **logic-based** agent programming language based on the **BDI** architecture and named it AgentSpeak(L).
- This became a **highly cited paper** in the multi-agent systems literature.
- AgentSpeak was an abstract agent programming language aimed to help the understanding of the relation between **practical implementations of the BDI architecture** such as PRS and the formalisation of the ideas behind the BDI architecture using **modal logics**.
- **Jason** is an interpreter for an extended version of AgentSpeak.
 - It implements the operational semantics of that language, and provides a platform for the development of multi-agent systems, with many user-customisable features.

Examples

Jason .mas2j project

```
MAS bdi_hw {  
    agents: bob;  
    // file bob.asl at project directory  
}
```

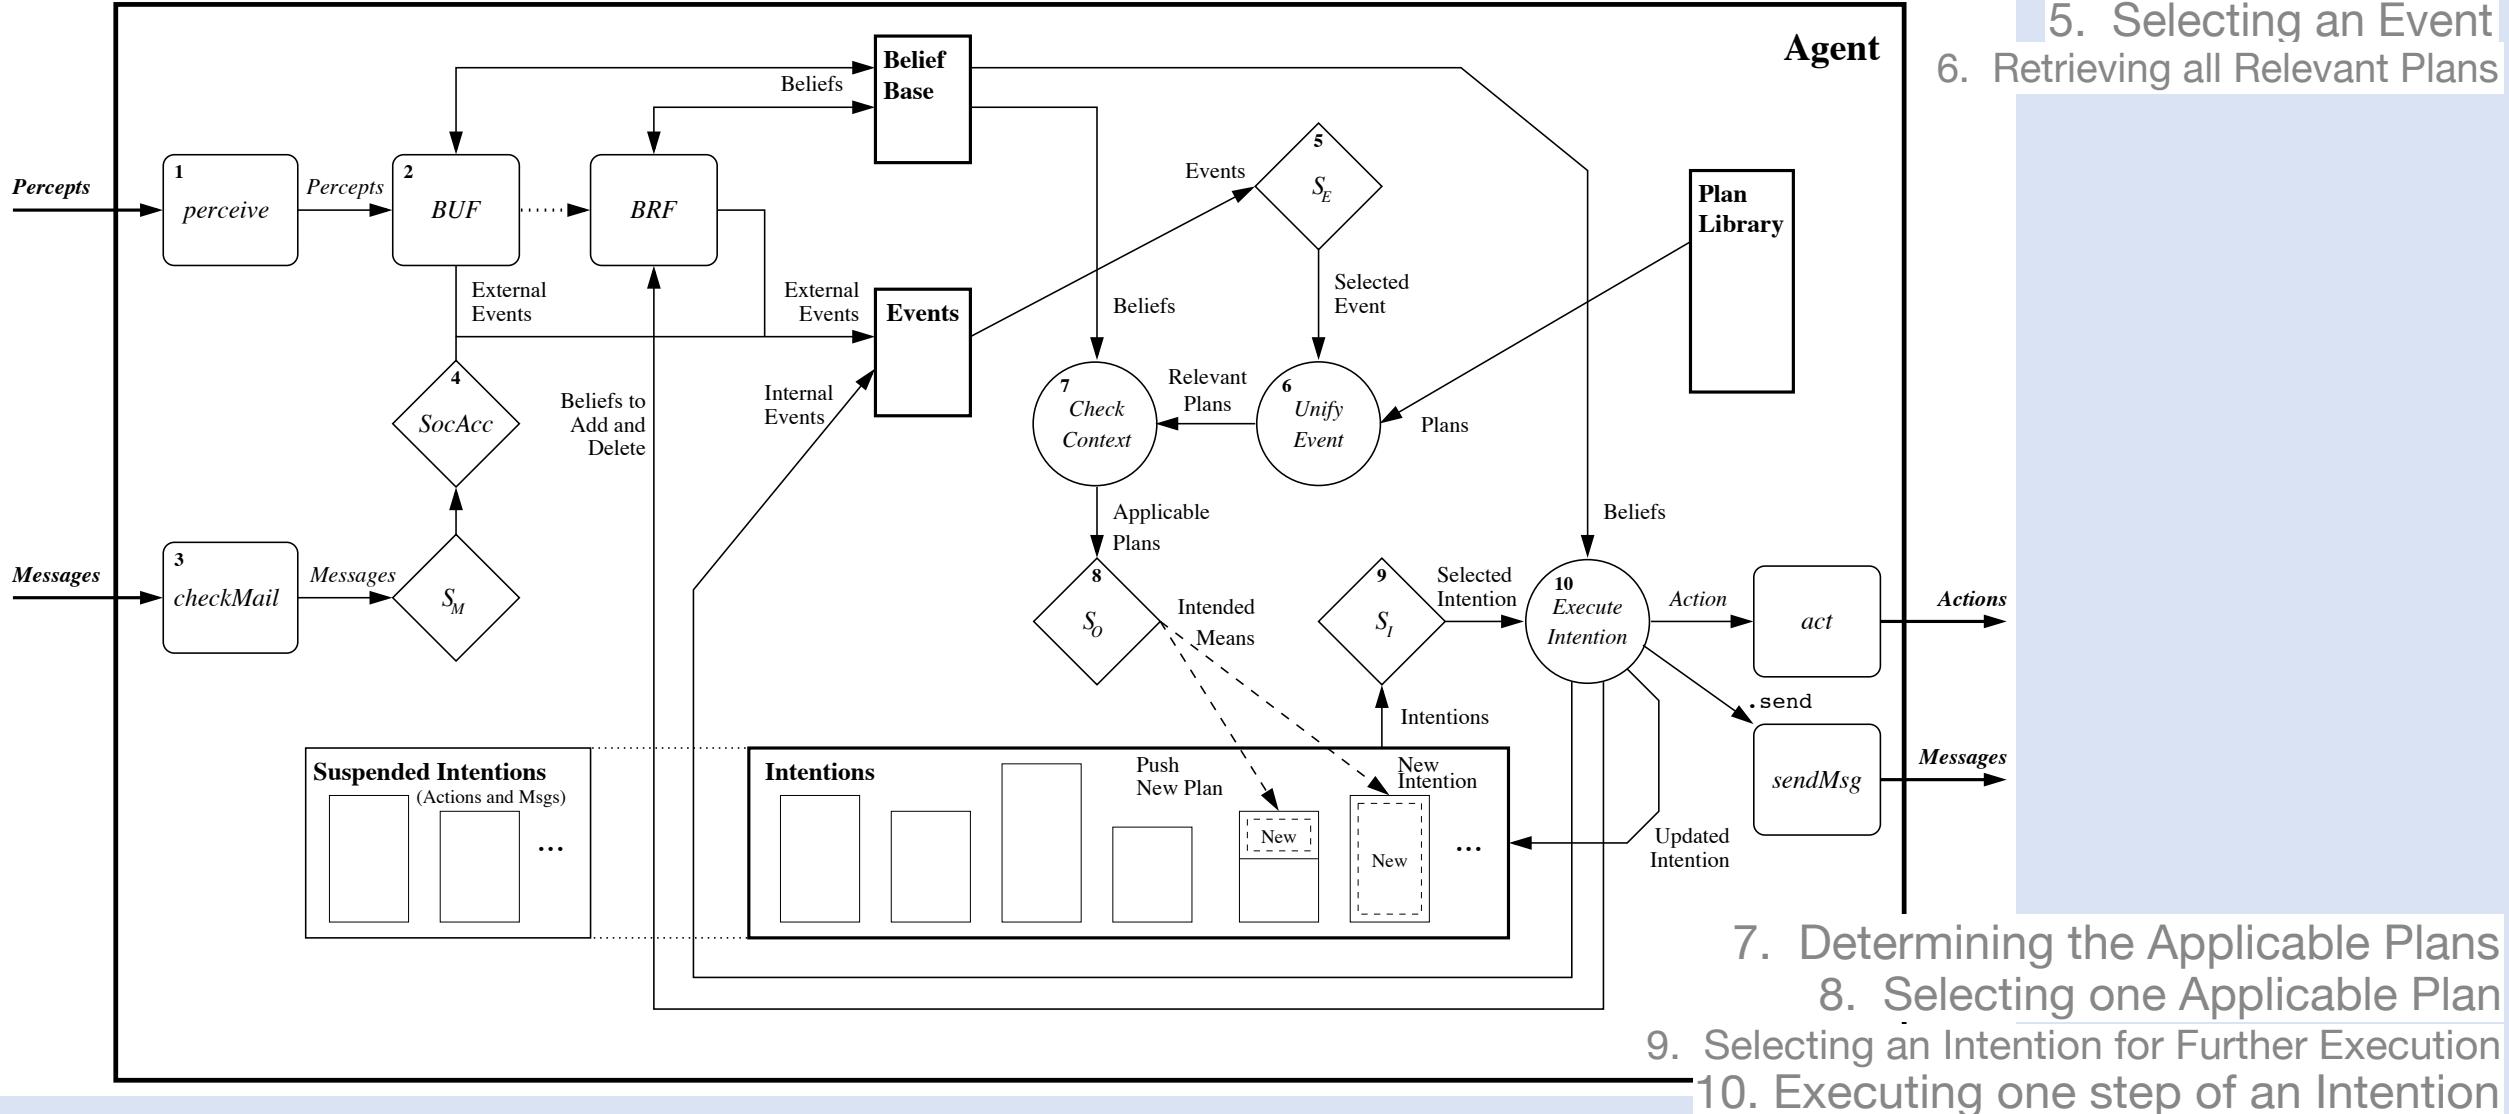
```
happy(bob).  
  
!say(hello).  
  
+!say(X) : happy(bob) <- .print(X).
```

```
+green_patch(Rock)  
: not battery_charge(low)  
<- ?location(Rock,Coordinates);  
    !at(Coordinates);  
    !examine(Rock).  
  
+!at(Coords)  
: not at(Coords)  
    & safe_path(Coords)  
<- move_towards(Coords);  
    !at(Coords).  
  
+!at(Coords) ...
```

Jason Reasoning Cycle

1. Perceiving the Environment
2. Updating the Belief Base
3. Receiving Communication from Other Agents
4. Selecting ‘Socially Acceptable’ Messages

5. Selecting an Event
6. Retrieving all Relevant Plans



7. Determining the Applicable Plans
8. Selecting one Applicable Plan
9. Selecting an Intention for Further Execution
10. Executing one step of an Intention

Agent Communication

- Agent communication in multi-agent systems is typically based on the **speech-act theory**
- Speech-act theory starts from the principle that language is action:
 - a rational agent makes an **utterance** in an attempt to change the state of the world, in the same way that an agent performs ‘physical’ actions to change the state of the world.
 - what distinguishes speech acts from other (‘non-speech’) actions is that the domain of a speech act –the part of the world that the agent wishes to modify through the performance of the act – is typically limited the mental state(s) of the hearer(s) of the utterance.
 - thus **examples of speech** acts might be to **change your beliefs, desires or intentions.**

Agent Communication

- The various types of speech acts are generally referred to as '**performatives**' in the context of agent communication.
 - representatives – such as informing, e.g. 'It is raining';
 - directives – attempts to get the hearer to do something, e.g. 'Please make the tea';
 - commisives – which commit the speaker to doing something, e.g. 'I promise to...';
 - expressives – whereby a speaker expresses a mental state, e.g. 'Thank you!';
 - declarations – such as declaring war or christening.
 - performative = request
content = 'the door is closed'
speech act = 'please close the door'
 - performative = inform
content = 'the door is closed'
speech act = 'the door is closed!'
 - performative = inquire
content = 'the door is closed'
speech act = 'is the door closed?'

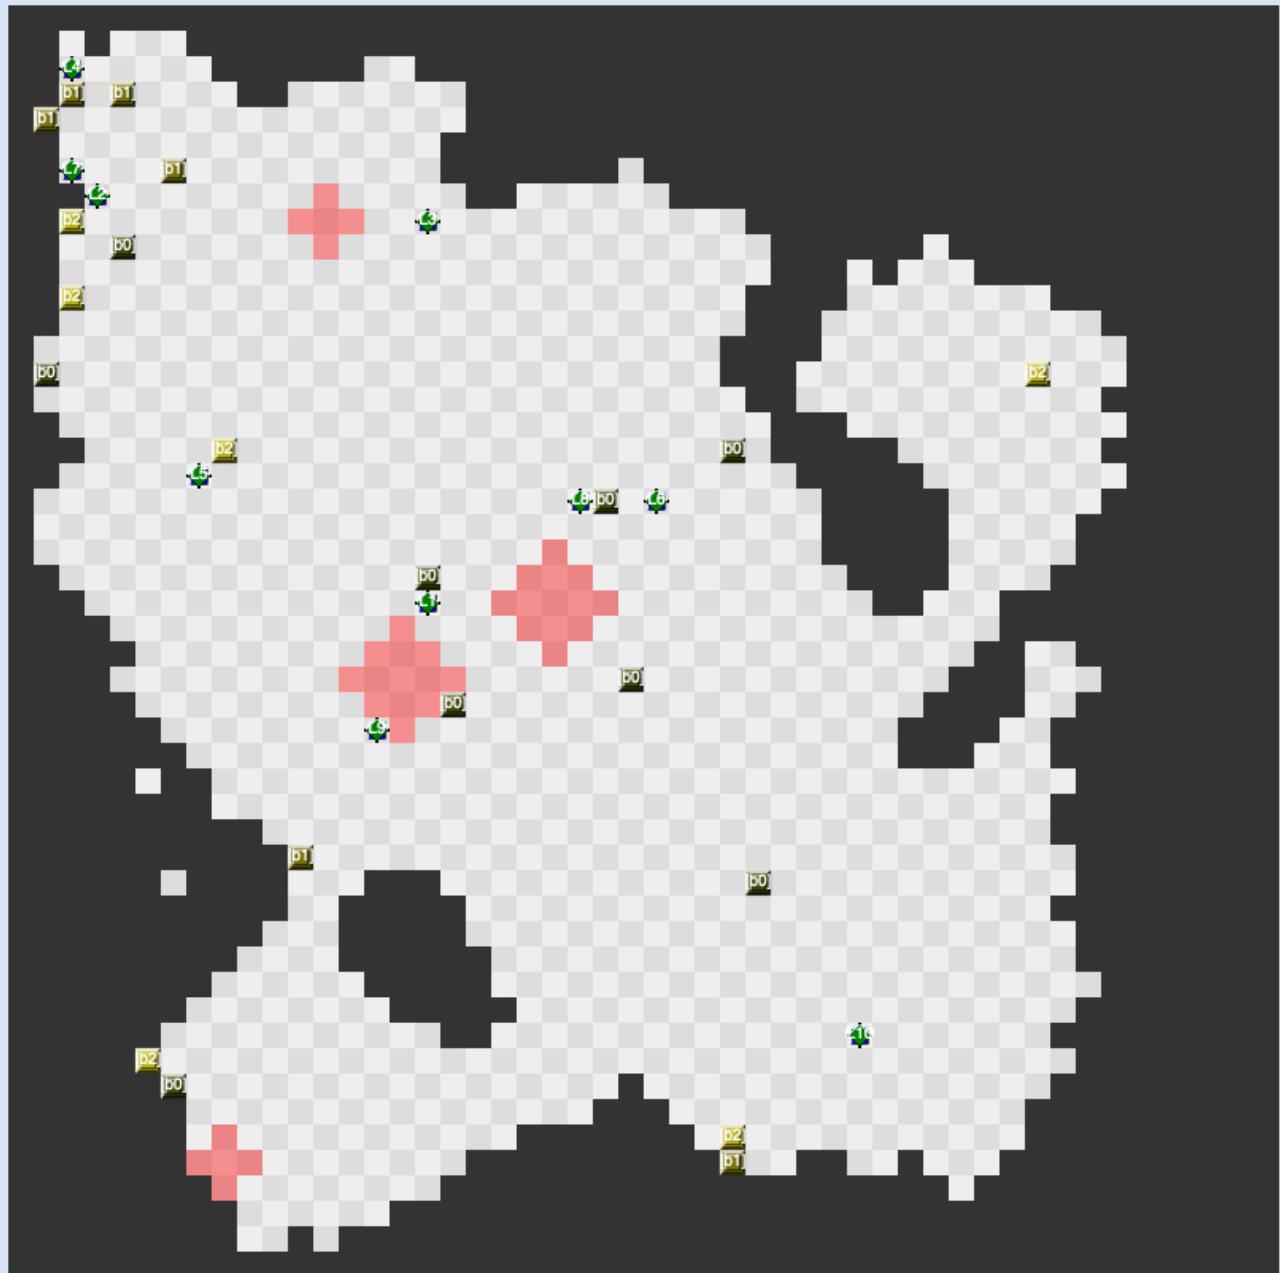
Agent Communication

- The Knowledge Query and Manipulation Language (**KQML**) was the first attempt to define a practical agent communication language that included high-level (speech-act based) communication

```
(ask-one
  :content (PRICE IBM ?price)
  :receiver stock-server
  :language LProlog
  :ontology NYSE-TICKS
)
```

- The **FIPA standard** for agent communication (**ACL**) was released in 2002.
 - The goal of the FIPA organisation was to develop a coherent collection of standards relating to agent communication, management and use, and the agent communication standard was just one component of the FIPA suite.

The Multi-Agent Programming Contest



The MASSim Platform

Server: Unified scenario implementation

Agents: Connect remotely (sockets)

- Receive percepts
(state of the visible environment)
- Deliberate and send actions back

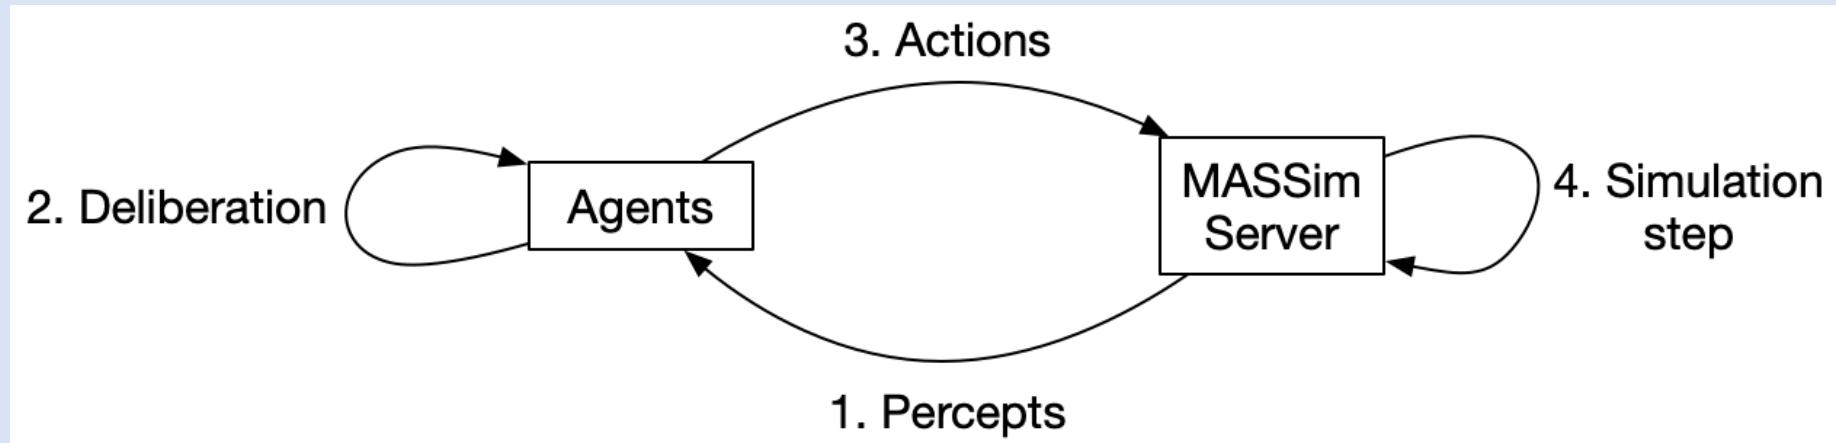
Simulation: Discrete steps

- Actions of all agents processed after each step

Tournament

- Games between **2 teams** of agents (as of now)
- Each team plays against each other
- Deliberation: **4 seconds** per step for network roundtrip and action computation
- **3 simulations** with different parameters per match
 - environment characteristics
 - agent characteristics (features, number, ...)

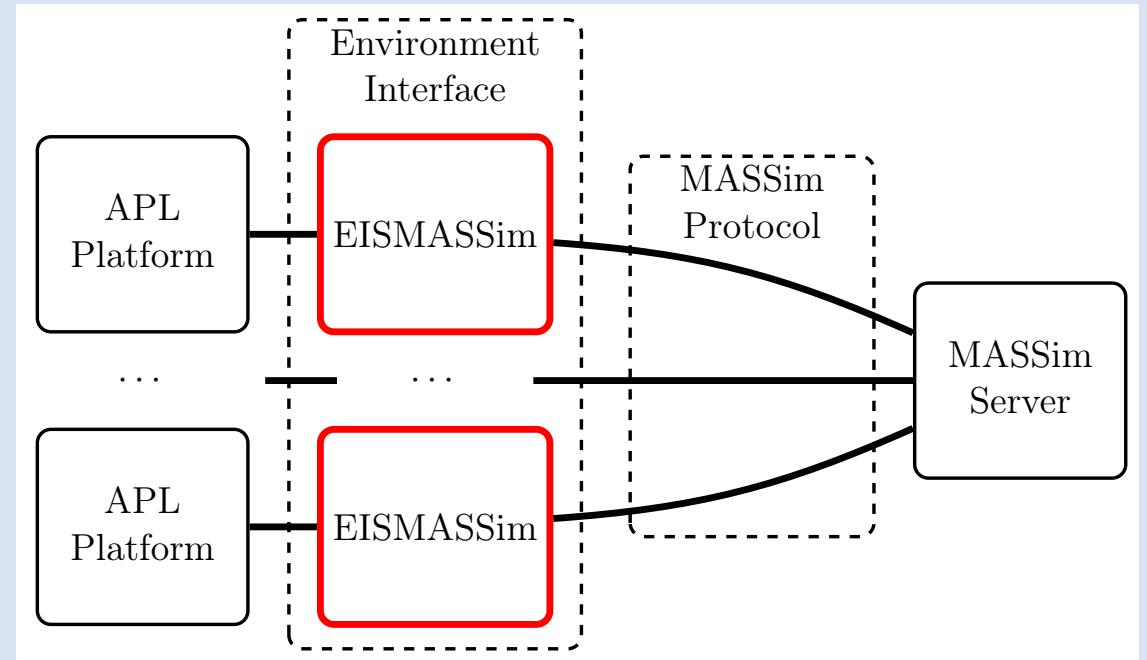
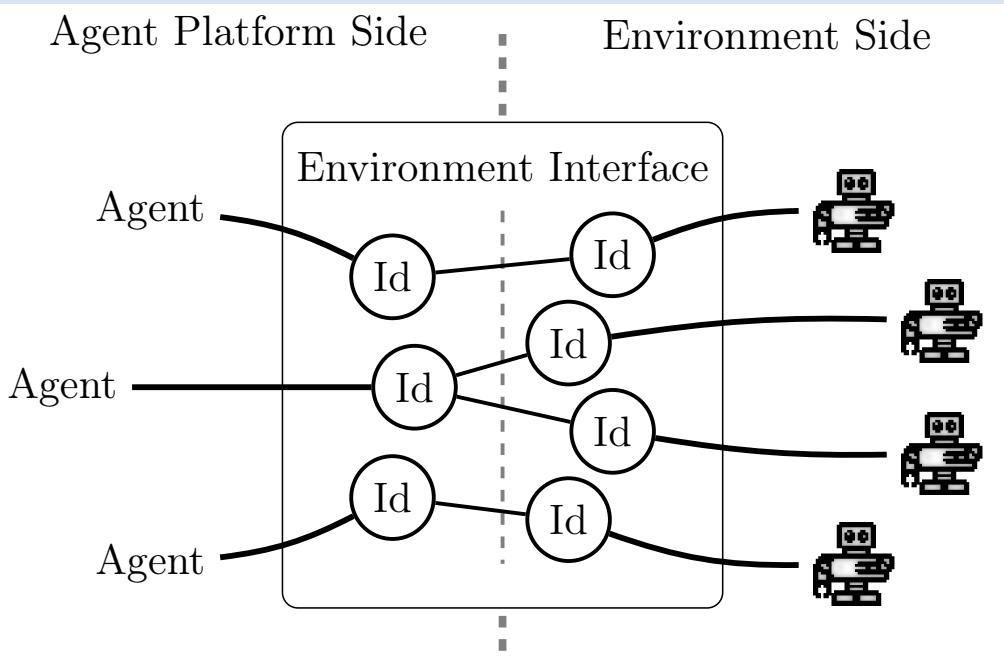
The MASSim Platform



Features in the MASSim

- The MASSim server is implemented in **Java**
- Highly **configurable**
 - Simulations
 - Steps
 - Teams
 - ...
- Also, it provides an implementation of EIS [2, 3] (EISMASSim) to Agent Platforms in Java
 - Abstraction to model the **agent-environment** interaction¹
 - Handles the **communication** with the MASSim server

Mapping Entities to Agents



Scenario



The End!

