

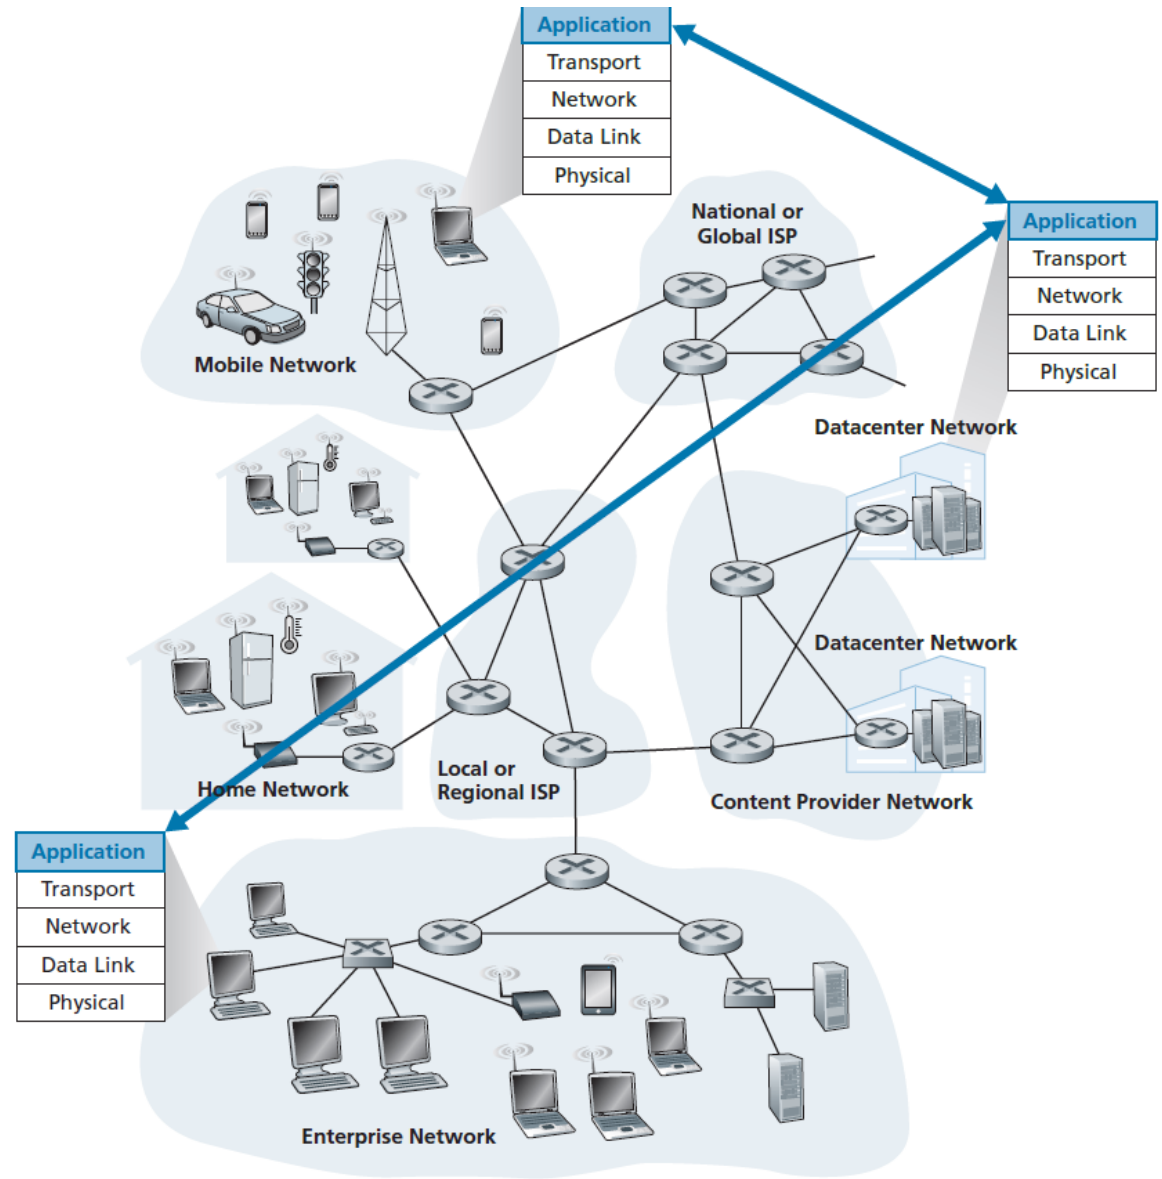
Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- Video streaming and content distribution networks
- Socket programming with UDP and TCP

Some network apps

need networking infrastructure and protocols (*since we use network applications*), which way to explain?

- social networking
- Web
- text messaging
- e-mail
- multi-user network games
- streaming stored video (YouTube, Netflix)
- P2P file sharing
- voice over IP (Skype)
- real-time video conferencing (Zoom)
- Internet search
- remote login
- ...



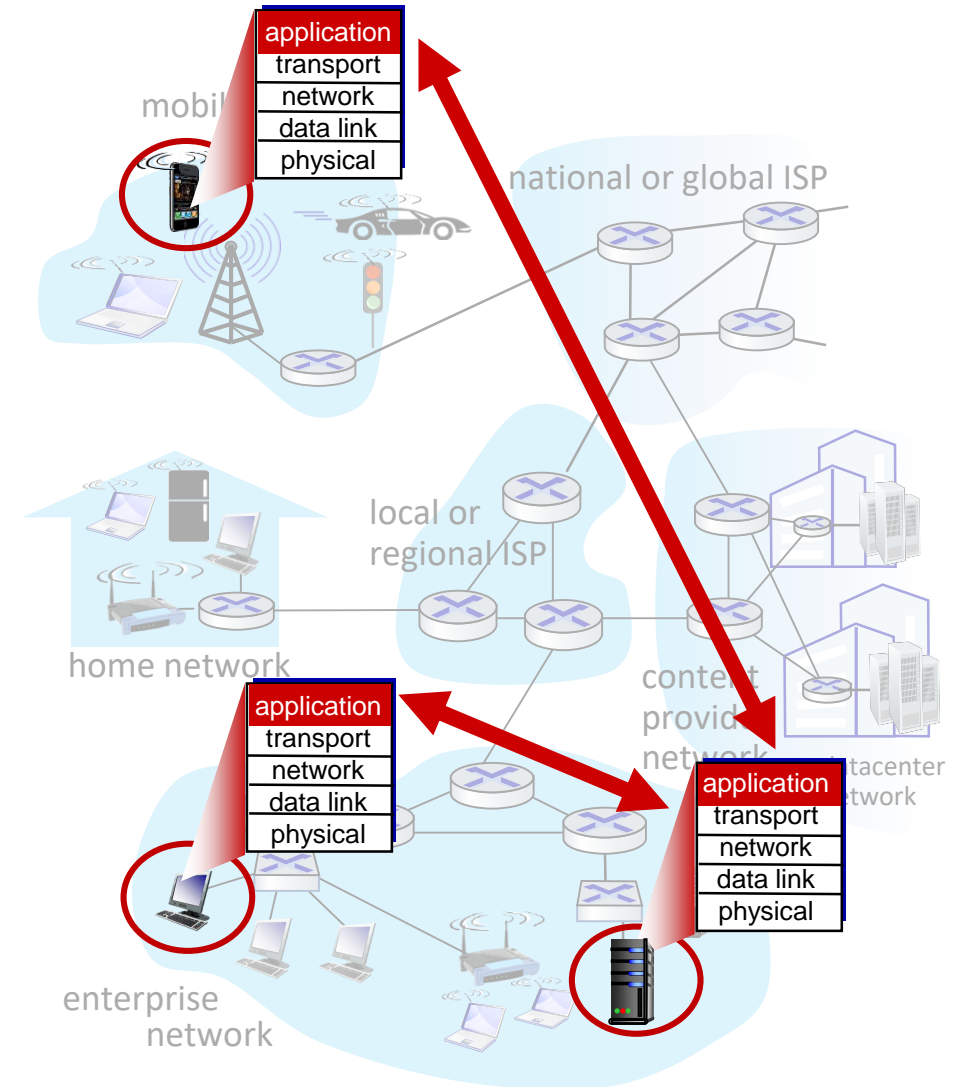
Creating a network app

writing programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software, Netflix-provided client program and and a Netflix server program

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



Creating a network app


Before starting coding, have an architectural plan for the application.

Application's architecture is different from the network architecture.

For an application developer, the network architecture is fixed (*e.g. five-layered model and the provided services via the protocols*).

Application architecture: designed by the developer, about how the application is structured over the end systems.

Two common types: client-server and peer-to-peer (P2P).




Client-server paradigm

server:

- always-on host
- permanent IP address
- often in data centers, for scaling
- need for multiple services, NLB

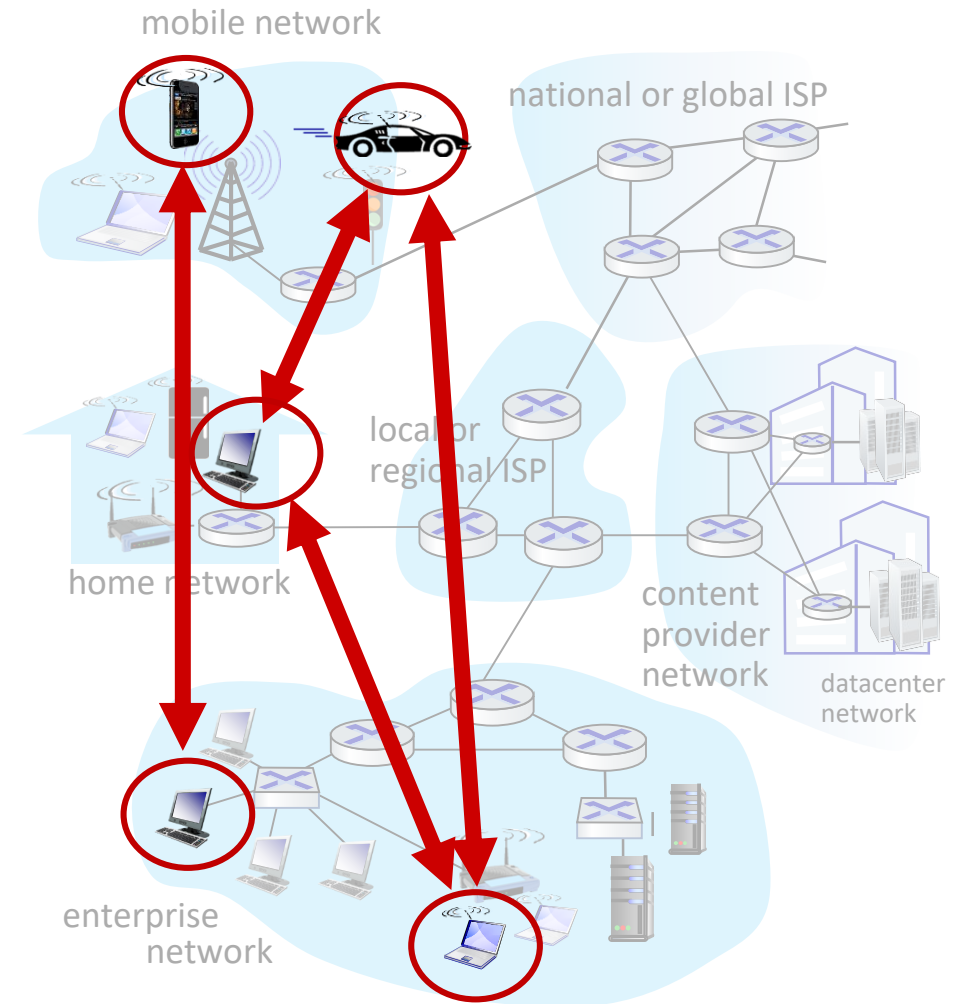
clients:

- contact, communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do *not* communicate directly with each other
- examples: HTTP, IMAP, FTP



Peer-peer architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management, security, performance, and reliability issues
 - scalability, cost-effectiveness
- example: P2P file sharing (*each peer adds service capacity to the overall system by distributing files to other peers*)



Processes communicating

Firstly, understand of how the programs, running in multiple end systems, communicate with each other.

For OS's, processes communicate.

Process: a program running in a host.

Different processes on a single host communicate via interprocess communication, rules are governed by OS.

Processes on different end systems communicate by exchanging messages across the network.

Processes communicating

process: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**

clients, servers

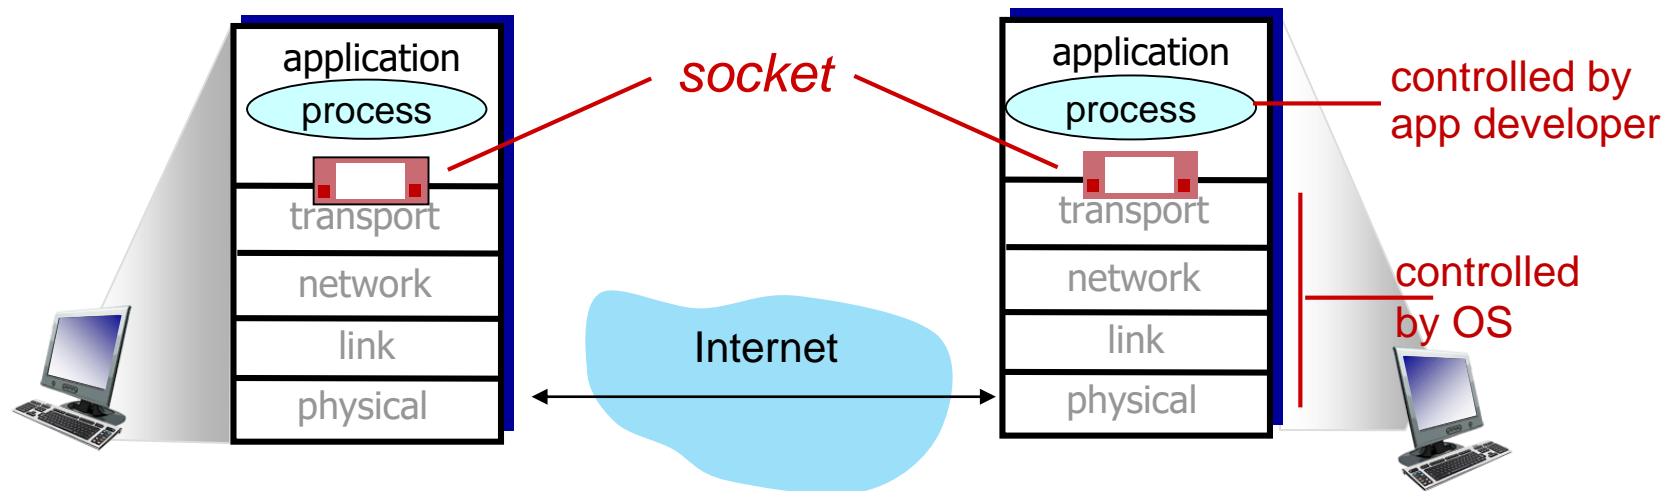
client process: process that initiates communication

server process: process that waits to be contacted

- note: applications with P2P architectures have client processes & server processes (*peers downloading and uoloading*)

Sockets

- process sends/receives messages to/from through its *socket (a software interface)*
- interface between the application layer and the transport layer within a host
- API
- developer controls application side, OS controls transport side
- process analogous to house, socket analogous to door
 - sending process shoves message out door
 - sending process relies on transport infrastructure
 - two sockets involved: one on each side



Addressing processes

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
- *identifier* includes both IP address and port numbers associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
 - IP address: 128.119.245.12
 - port number: 80

An application-layer protocol defines:

- **types of messages exchanged**,
 - e.g., request, response
- **message syntax**:
 - what fields in messages & how fields are delineated
- **message semantics**
 - meaning of information in fields
- **rules** for when and how processes send & respond to messages

open protocols:

- defined in RFCs, everyone has access to protocol definition
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype, Zoom

What transport service does an app need?

More than one protocols, how to evaluate and pick one?

reliability (correctly and completely delivery)

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio, multimedia) can tolerate some loss – loss tolerant

timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

throughput: a guaranteed available throughput, fixed rate

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective” – bandwidth sensitive
- other apps (“elastic apps”) make use of whatever throughput they get

security

- encryption, data integrity, ...

Transport service requirements: common apps

TCP or UDP?

application	data loss	throughput	time sensitive?
file transfer/download	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5Kbps-1Mbps video:10Kbps-5Mbps	yes, 10's msec
streaming audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	Kbps+	yes, 10's msec
text messaging	no loss	elastic	yes and no

Internet transport protocols services

TCP service:

- *reliable transport* between sending and receiving process (*without error and in the proper order*)
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *connection-oriented*: setup required between client and server processes
- *does not provide*: timing, minimum throughput guarantee, security

UDP service:

- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup.

Internet applications, and transport protocols

application	application layer protocol	transport protocol
-------------	----------------------------	--------------------

file transfer/download	FTP [RFC 959]	TCP
e-mail	SMTP [RFC 5321]	TCP
Web documents	HTTP 1.1 [RFC 7320]	TCP
Internet telephony	SIP – session initiation [RFC 3261], RTP – real time transport [RFC 3550], or proprietary	TCP or UDP
streaming audio/video	HTTP [RFC 7320], DASH	TCP
interactive games	WOW, FPS (proprietary)	UDP or TCP

Securing TCP

TCP does not provide security by itself, can be enhanced at the app. layer

Vanilla TCP & UDP sockets:

- no encryption
- cleartext passwords sent into socket traverse Internet in clear text

Transport Layer Security (TLS)

- provides encrypted TCP connections
- data integrity
- end-point authentication

TLS implemented in application layer

- apps use TLS libraries, that use TCP in turn
- cleartext sent into “socket” traverse Internet *encrypted*

Application layer: overview

- Principles of network applications
- **Web and HTTP**
- E-mail, SMTP, IMAP
- The Domain Name System
DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP

Web and HTTP

- an application-layer protocol defines ‘how an application’s processes, running on different end systems, pass messages to each other’
- HTTP is an application layer protocol, web is a client-server application (network app.)
- defines the format and sequence of messages (between browser and Web server)
- implemented in two programs: client and a server
- web page (document) consists of *objects*, each of which can be stored on different Web servers
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects, each* addressable by a *URL*, e.g.,

`www.someschool.edu/someDept/pic.gif`


host name

path name

HTTP overview

HTTP: hypertext transfer protocol

- Web's application-layer protocol
- client/server model:
 - *client*: browser that requests, receives, (using HTTP protocol) and “displays” Web objects
 - *server*: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

HTTP uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is “stateless”

- server maintains *no* information about past client requests (if client asks for the same object twice, server resends the object)

aside
protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

HTTP connections: two types

Non-persistent HTTP

1. TCP connection opened
 2. at most one object sent over TCP connection
 3. TCP connection closed
- downloading multiple objects required multiple connections

Persistent HTTP

- TCP connection opened to a server
- multiple objects can be sent over *single* TCP connection between client, and that server
- TCP connection closed

Non-persistent HTTP: example

User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)



1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80



1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80 “accepts” connection, notifying client

2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

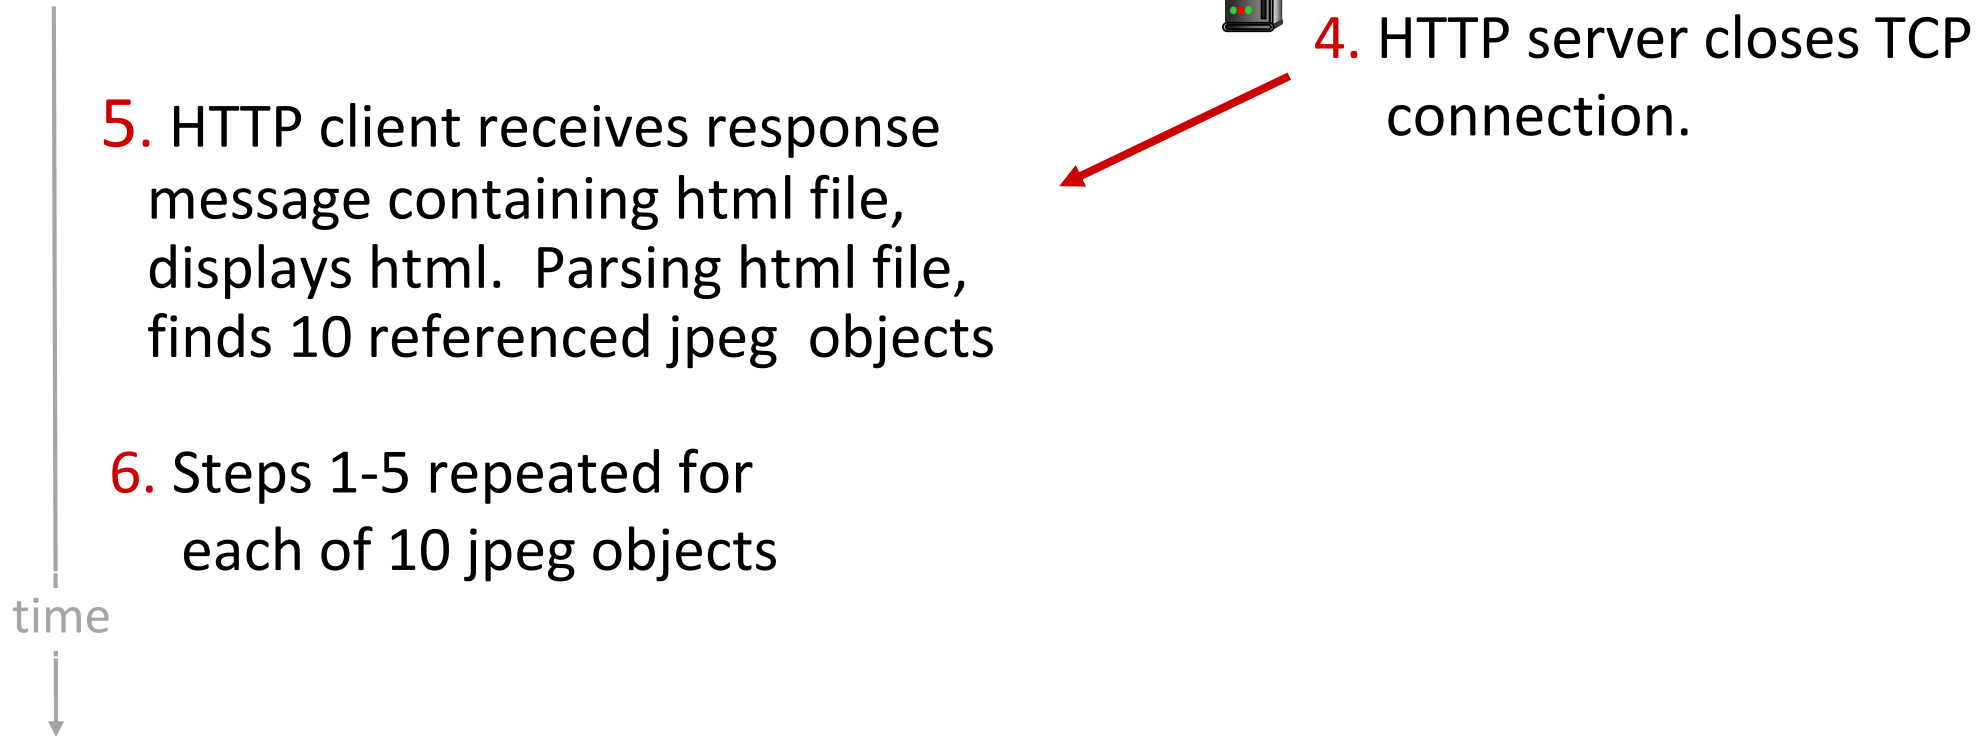
3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time



Non-persistent HTTP: example (cont.)

User enters URL: `www.someSchool.edu/someDepartment/home.index`
(containing text, references to 10 jpeg images)

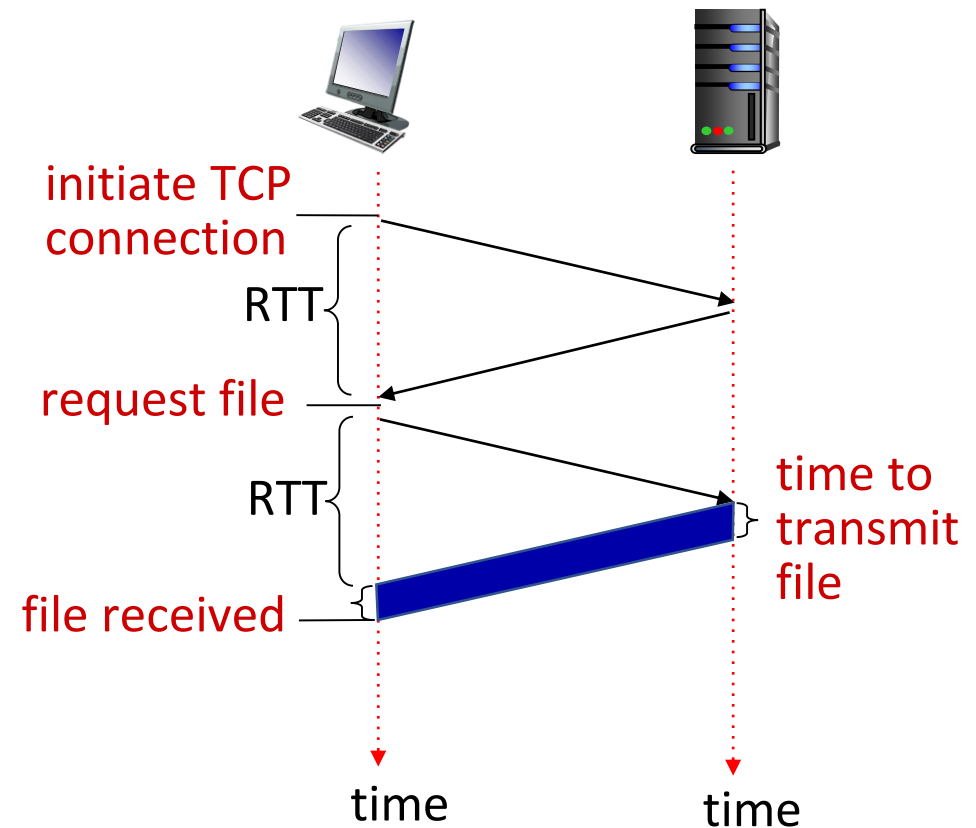


Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back, includes packet delays

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



Non-persistent HTTP response time = $2RTT + \text{file transmission time}$

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel

Persistent HTTP (HTTP1.1): default one

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

HTTP request message

- two types of HTTP messages: *request, response*

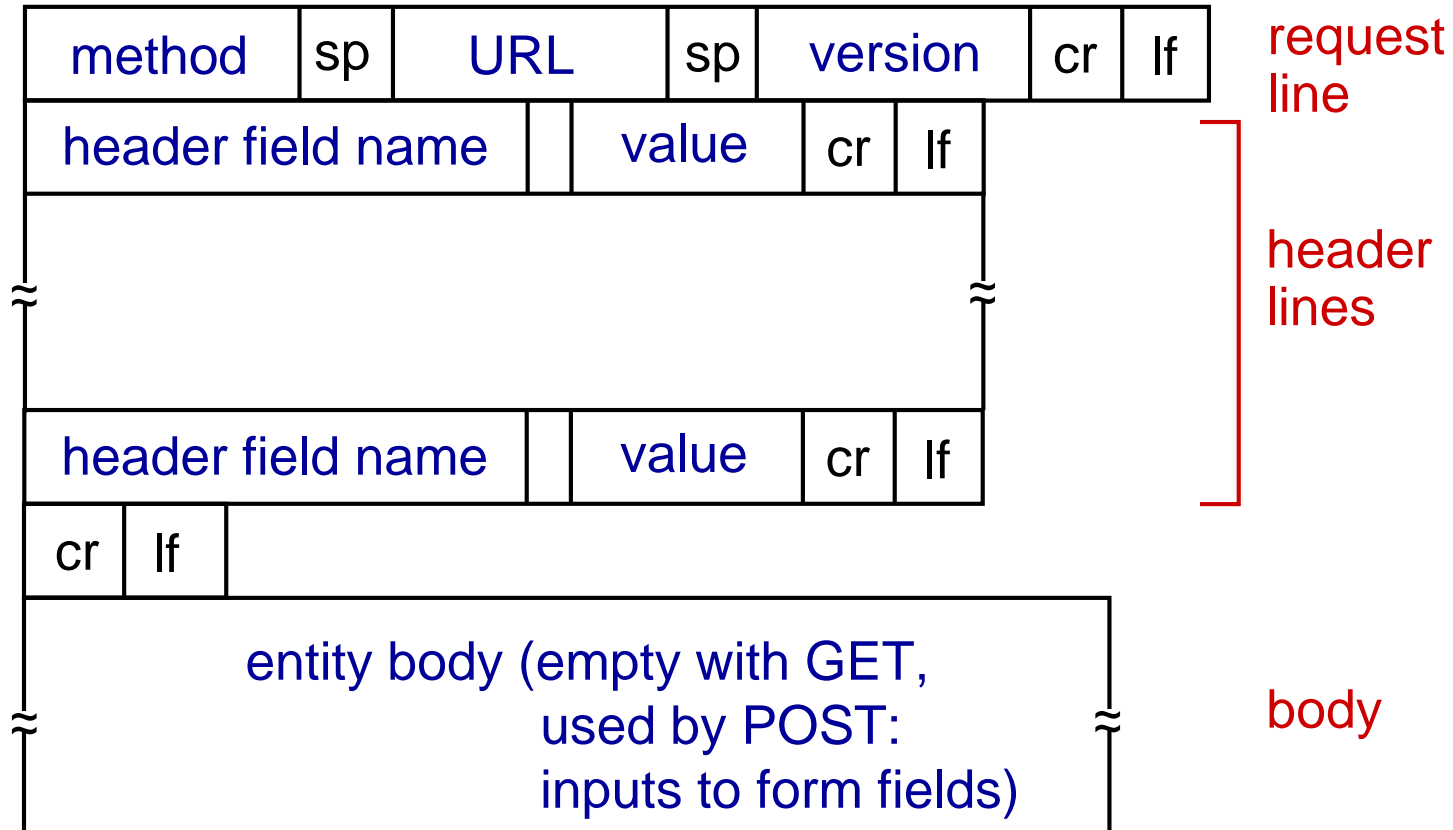
- **HTTP request message:**

- ASCII (human-readable format)

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
Connection: close
User-agent: Mozilla/5.0
Accept-language: fr
```

- first line: req. line, three fields: method, URL, and HTTP version
- method: **GET**, POST, HEAD, PUT, or DELETE
- GET: requested object is identified in URL
- subsequent lines: header lines
 - host: web server
 - conn. close: 'close of conn. after sending the requested object' claimed from the server
 - UA: browser type
 - Lang.: preferred version of the object requested

HTTP request message: general format



Other HTTP request messages

POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a '?'):

www.somesite.com/search?name&telephoneNumber

HEAD method:

- for debugging
- server responds with an HTTP message but it leaves out the requested object

PUT method:

- uploads new file (object) to server

DELETE method:

- deletes an object on a server

HTTP response message

```
HTTP/1.1 200 OK
Connection: close
Date: Tue, 18 Aug 2015 15:44:04 GMT
Server: Apache/2.2.3 (CentOS)
Last-Modified: Tue, 18 Aug 2015 15:11:03 GMT
Content-Length: 6821
Content-Type: text/html

(data data data data data ...)
```

- three sections: an initial status line, header lines, and then the entity body
- entity body contains the requested object
- status line has three fields: the protocol version field, a status code, and a corresponding status message
- Connection: close header line: server will close the TCP connection after sending
- Date: header line: the time and date when the HTTP response was created and sent
- Server: header line: the Web server
- Last-Modified: header line: the time and date when the object was created (or last modified)
- Content-Length: header line: the number of bytes in the object
- Content-Type: header line: object in the entity body is HTML text

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field), client software will automatically retrieve the new URL

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

Maintaining user/server state: cookies

(since HTTP is stateless)

Web sites and client browser use *cookies* to maintain some state between transactions, sites can keep track of users

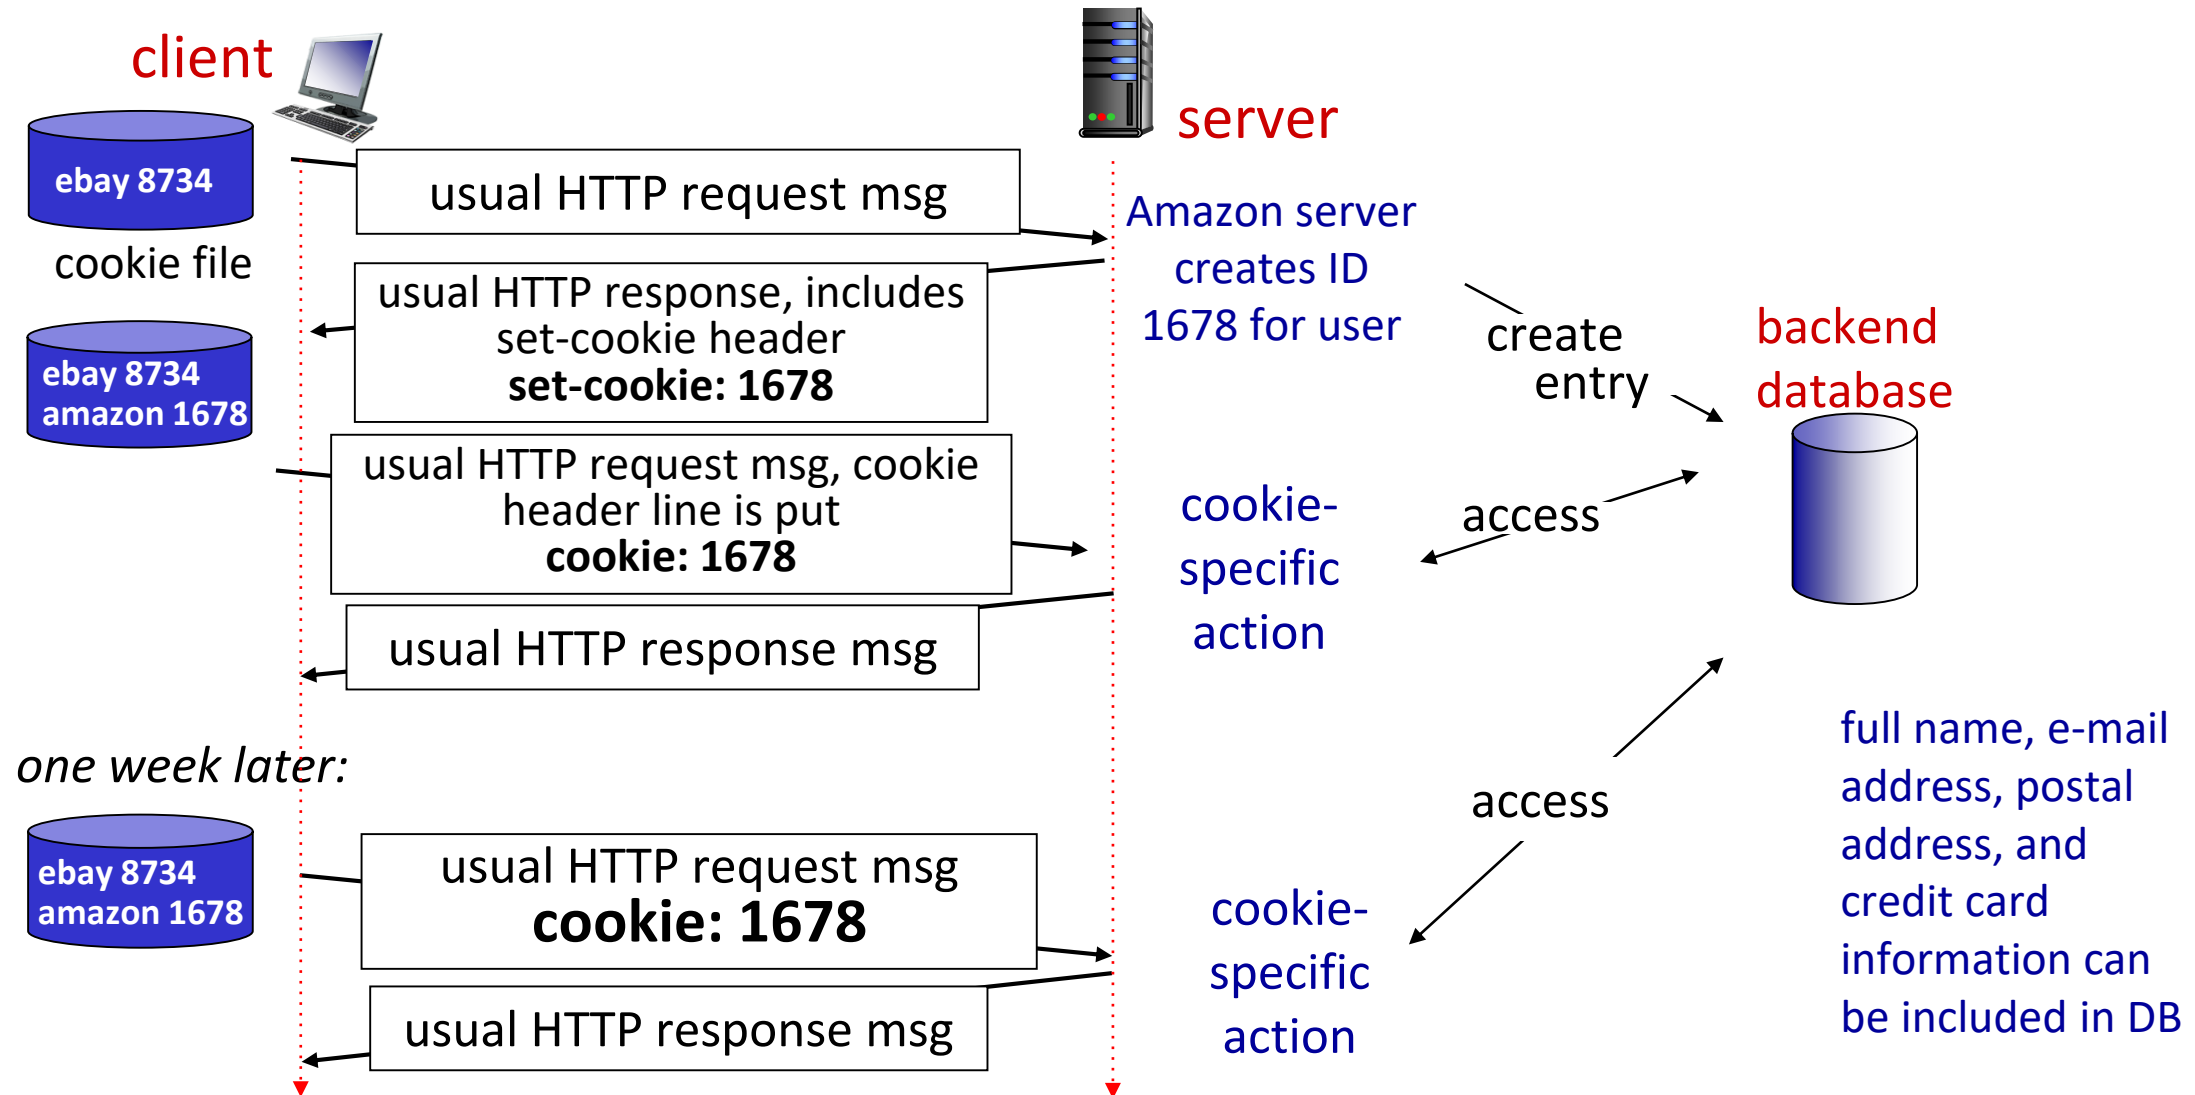
four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID (aka “cookie”)
 - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to “identify” Susan


Maintaining user/server state: cookies



Web caches (proxy)

Goal: satisfy client requests without involving origin server, respond to HTTP requests on the behalf of a Web server

- receives the object, stores a copy in its local storage and sends a copy to the client
- user configures browser to point to a (local) *Web cache*
- browser sends all HTTP requests to cache
 - *if* object in cache: cache returns object to client
 - *else* cache requests object from origin server, caches received object, then returns object to client



HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client.

- methods, status codes, most header fields unchanged from HTTP 1.1
- changes how the data is formatted and transported
- HTTP/1.1 uses persistent TCP connections (by default)

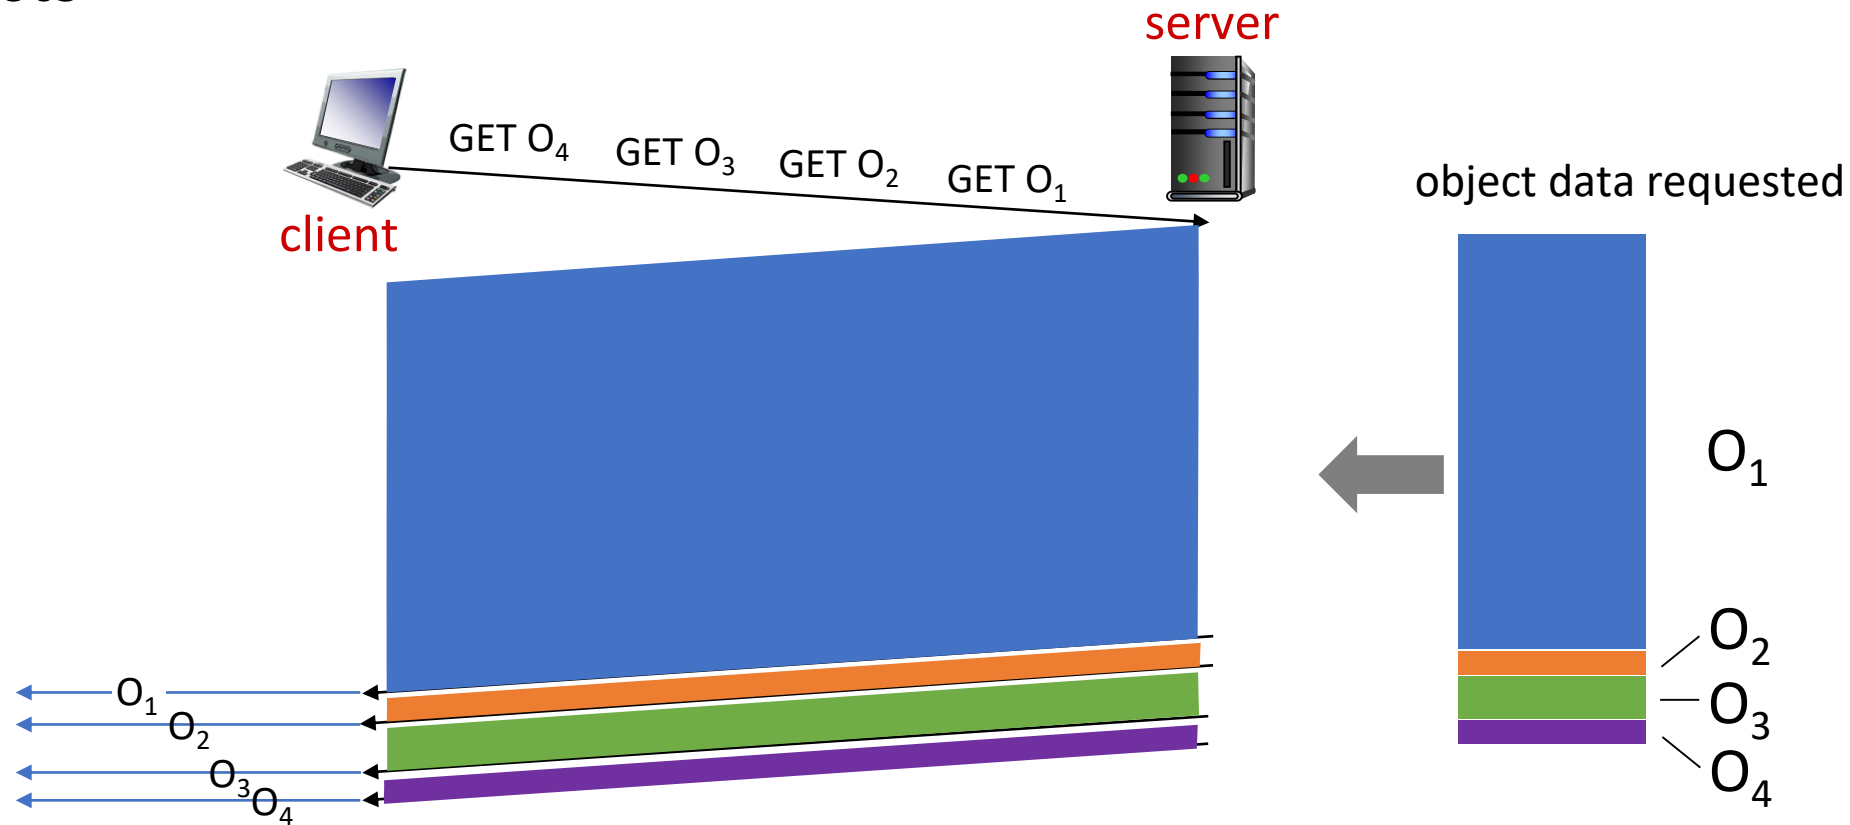
a Web page is sent from server to client over a single TCP conn.

number of sockets at the server is reduced

has a Head of Line (HOL) blocking problem

HTTP/2: mitigating HOL blocking

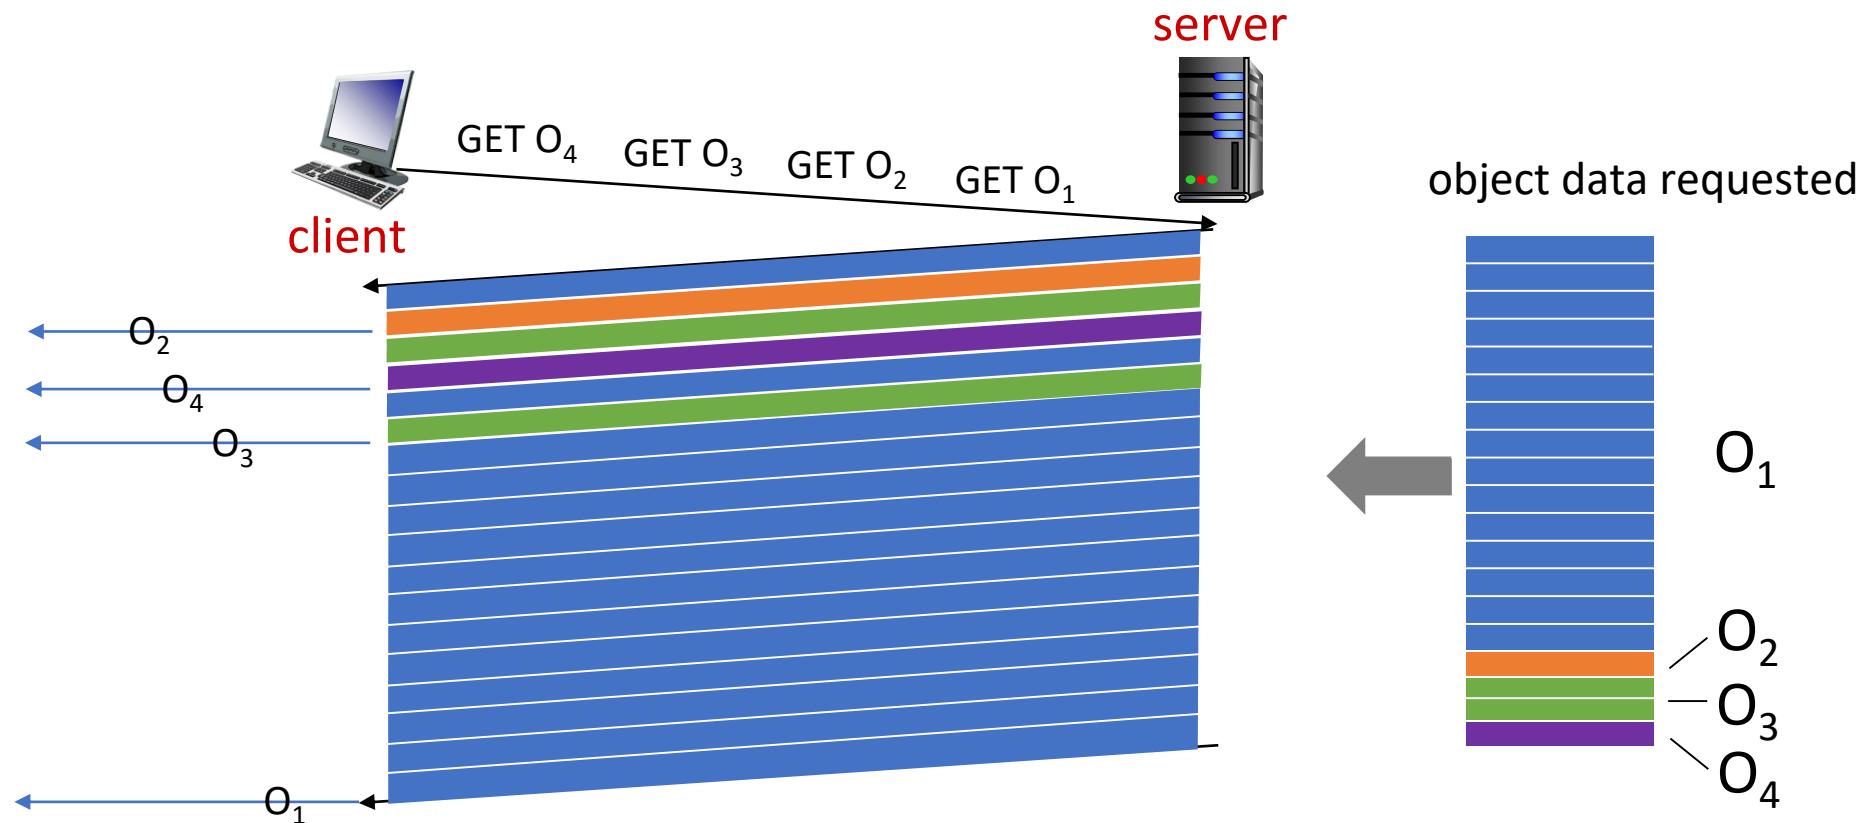
HTTP 1.1: client requests one large object (e.g., video file) and 3 smaller objects



objects delivered in order requested: O_2 , O_3 , O_4 wait behind O_1 ; can solve this problem by opening multiple parallel TCP connections – a good idea?

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission is interleaved
after sending one frame from the video clip, the first frames of each of the small objects are sent



O₂, O₃, O₄ delivered quickly, O₁ slightly delayed

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System
DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP


E-mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

User Agent

- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, iPhone mail client
- outgoing, incoming messages stored on server, user agent retrieves the message from mailboxes in a server



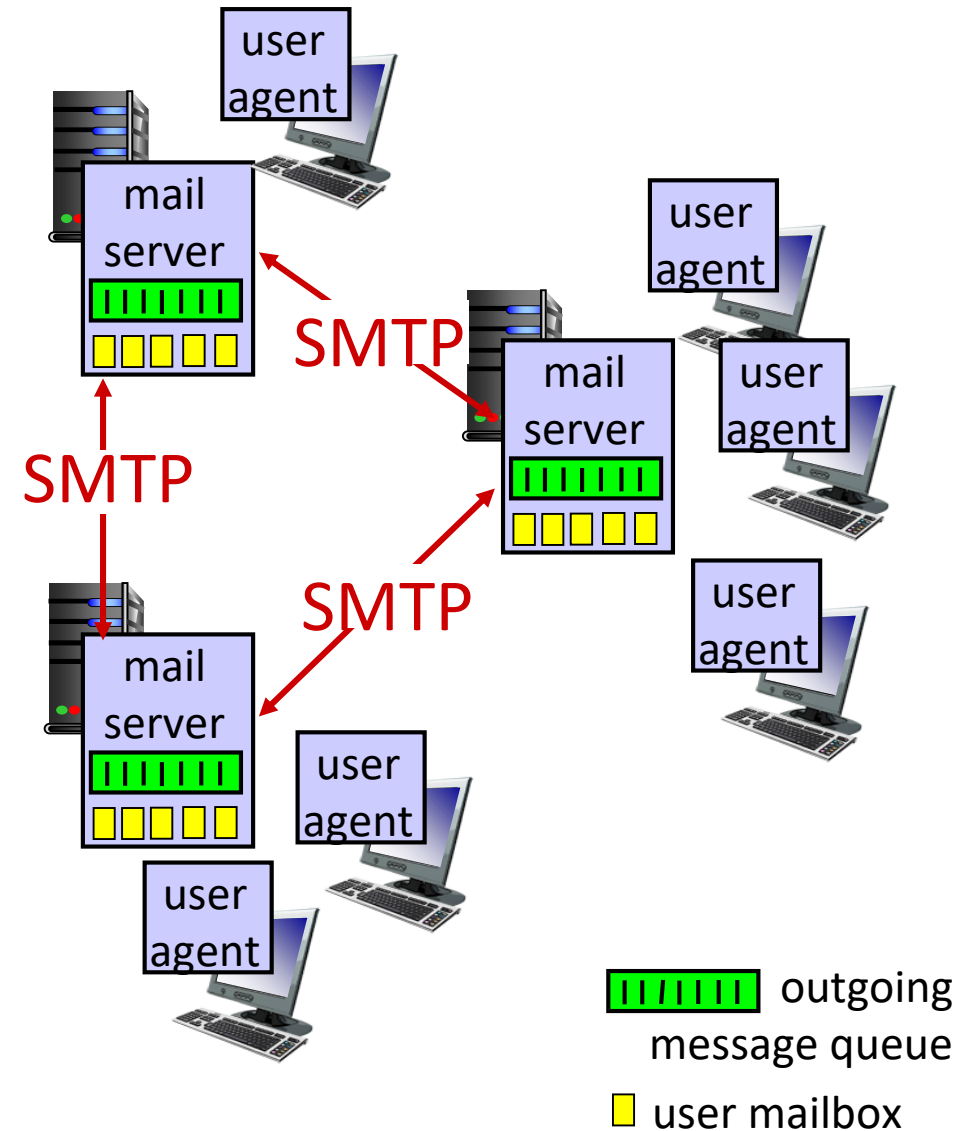
E-mail: mail servers

mail servers:

- *mailbox* contains incoming messages for user
- *message queue* of outgoing (to be sent) mail messages
- Senders mail server also deal with failures in recipients mail server (If the message cannot be delivered, senders server holds the message in a message queue and attempts to transfer the message later, reattempts are done every 30 mins)

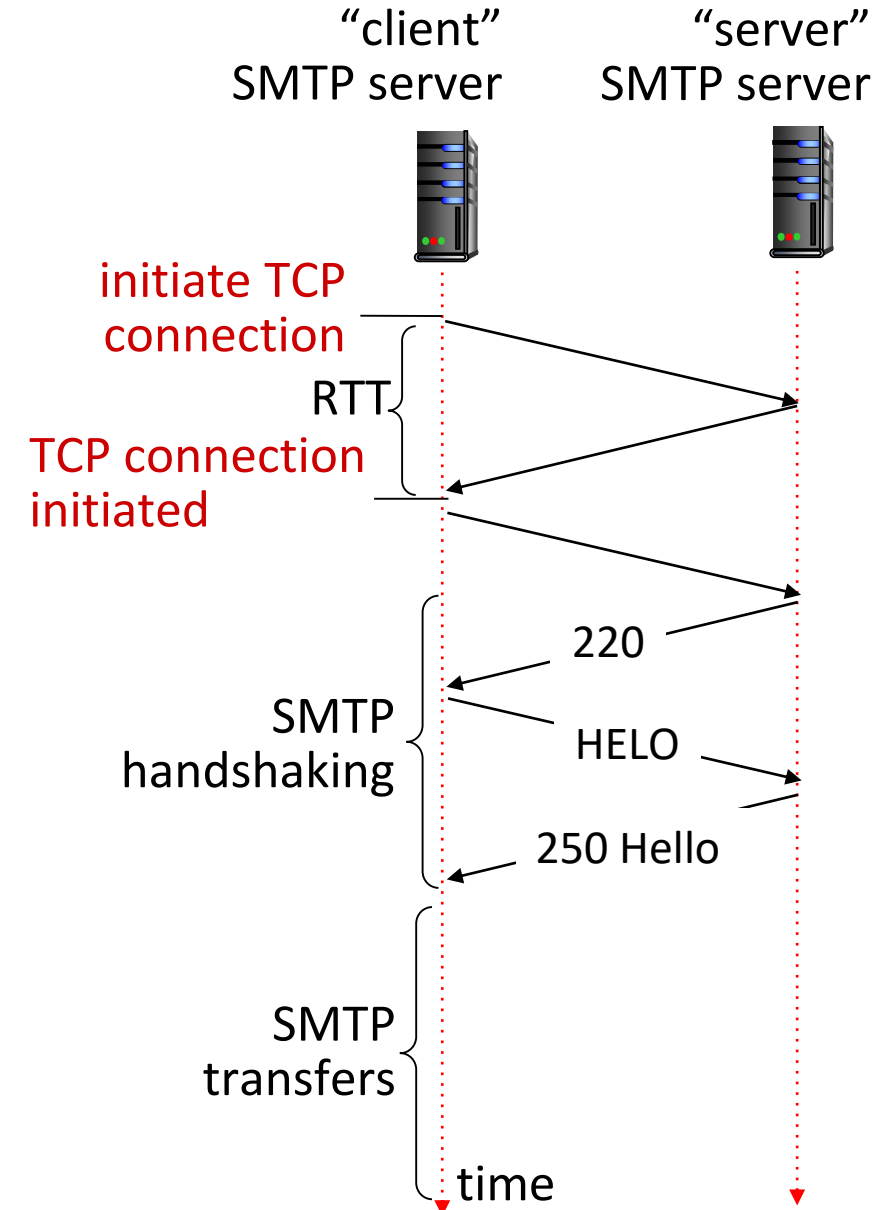
SMTP protocol between mail servers to send email messages

- **client**: sending mail server
- **“server”**: receiving mail server




SMTP RFC (5321)

- uses TCP to reliably transfer email message from client (mail server initiating connection) to server, port 25
 - direct transfer: sending server (acting like client) to receiving server
- three phases of transfer
 - SMTP handshaking (greeting)
 - SMTP transfer of messages
 - SMTP closure
- command/response interaction (like HTTP)
 - **commands:** ASCII text
 - **response:** status code and phrase



Scenario: Alice sends e-mail to Bob

- 1) Alice uses UA to compose e-mail message “to” bob@some school.edu
- 2) Alice’s UA sends message to her mail server using SMTP; message placed in message queue
- 3) client side of SMTP at mail server opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



Sample SMTP interaction


```
S: 220 hamburger.edu           (service ready)
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
(requested mail action okay)
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
(start mail input)
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
```

Mail message format

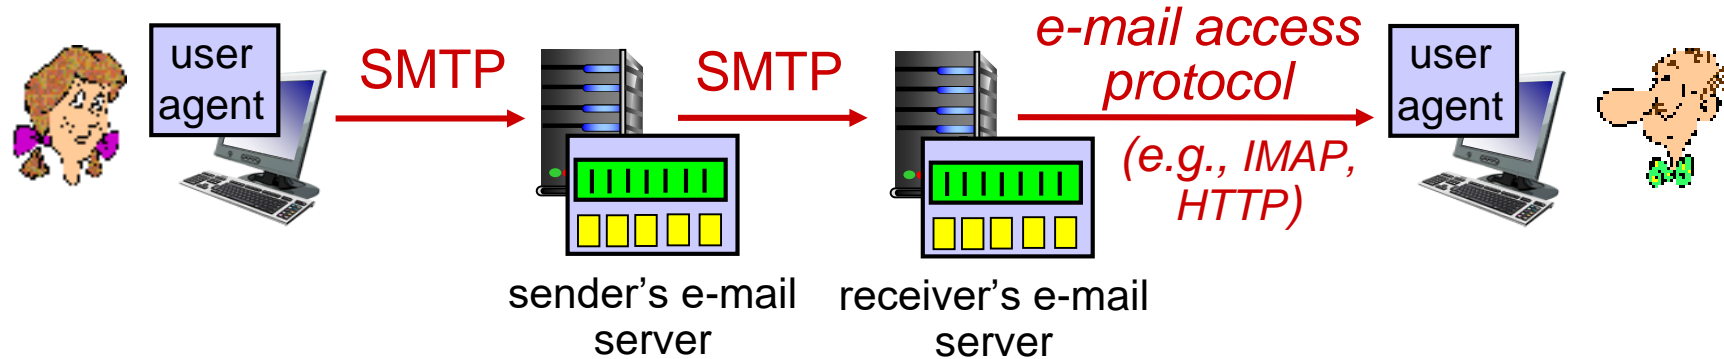
SMTP: protocol for exchanging e-mail messages, defined in RFC 5321 (like RFC 7231 defines HTTP)

RFC 2822 defines *syntax* for e-mail message itself (like HTML defines syntax for web documents)

- header lines, e.g.,
 - To:
 - From:
 - Subject:these lines, within the body of the email message area different from SMTP MAIL FROM:, RCPT TO: commands!
- Body: the “message” , ASCII characters only



Retrieving email: mail access protocols



- **SMTP:** delivery/storage of e-mail messages to receiver's server
- mail access protocol: retrieval from server
 - **IMAP:** Internet Mail Access Protocol [RFC 3501]: messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
- **HTTP:** gmail, Hotmail, Yahoo!Mail, etc. provides web-based interface on top of SMTP (to send), IMAP (or POP) to retrieve e-mail messages